#### ECE408/CS483/CSE408 Spring 2025 Applied Parallel Programming

# Lecture 6: Generalized Tiling & DRAM Bandwidth

#### Course Reminders

- Lab 2 is due this Friday
  - Make sure to implement all required code, including host memory allocation/deallocation.
  - Do not forget to push your work to github!
  - Submit your Lab 2 quiz only after you are done with the code.
- Lab 3 will be released soon

#### Getting help from our course chatbot

- Try asking these questions:
  - How do I allocate device memory?
  - How do I copy memory from host to device?
  - How do I call a GPU kernel?
  - Show an example of declaring a shared memory array in CUDA kernel.
  - Explain vectorAdd kernel code provided in the lectures.
  - Explain following code line by line in detail: "code goes here"
  - I am preparing for the exam, please help me to review the course materials from lectures 2 and 3. Ask me questions about it. Ask one question at a time, wait for my answer, and then check it for correctness.

#### Objectives

- To learn to handle boundary conditions in tiled algorithms.
- To understand the organization of memory based on dynamic RAM (DRAM).
- To understand the use of burst mode and multiple banks (both sources of parallelism) to increase DRAM performance (data rate).
- To understand memory access coalescing, which connects GPU kernel performance to DRAM organization.

#### How to Handle Matrices of Other Sizes?

- Lecture 5's tiled kernel
  - assumed integral number of tiles (thread blocks)
  - in all matrix dimensions.

How can we avoid this assumption?

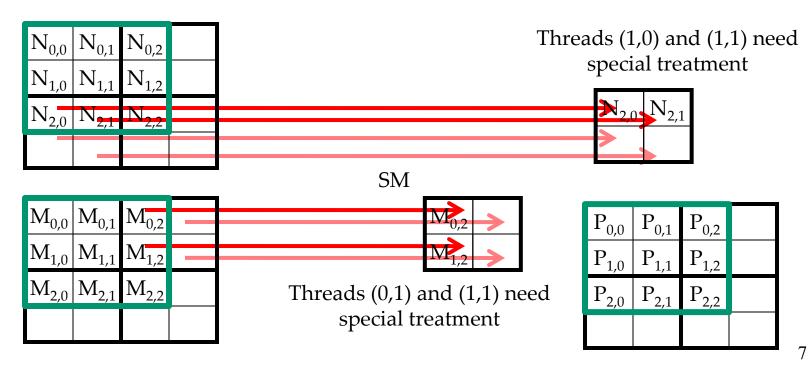
 One answer: add padding, but not easy to reformat data, and adds transfer time.

Other ideas?

#### Let's Review Our Kernel

```
global void MatrixMulKernel(float* M, float* N, float* P, int Width)
     shared float subTileM[TILE WIDTH] [TILE WIDTH];
     shared float subTileN[TILE WIDTH] [TILE WIDTH];
   int bx = blockIdx.x; int by = blockIdx.y;
4. int tx = threadIdx.x; int ty = threadIdx.y;
   // Identify the row and column of the P element to work on
5. int Row = by * TILE WIDTH + ty; // note: blockDim.x == TILE WIDTH
   int Col = bx * TILE WIDTH + tx; // blockDim.y == TILE WIDTH
   float Pvalue = 0;
   // Loop over the M and N tiles required to compute the P element
   // The code assumes that the Width is a multiple of TILE WIDTH!
   for (int m = 0; m < Width/TILE WIDTH; ++m) {
      // Collaborative loading of M and N tiles into shared memory
      subTileM[ty][tx] = M[Row*Width + m*TILE WIDTH+tx];
9.
      subTileN[ty][tx] = N[(m*TILE WIDTH+ty)*Width+Col];
10.
11.
      syncthreads();
12.
      for (int k = 0; k < TILE WIDTH; ++k)
13.
          Pvalue += subTileM[ty][k] * subTileN[k][tx];
14.
       syncthreads();
15. }
16. P[Row*Width+Col] = Pvalue;
```

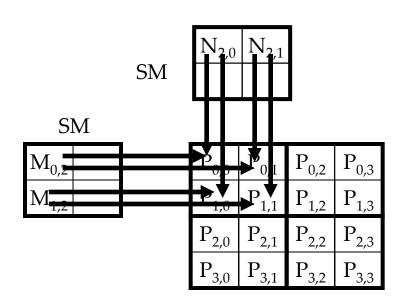
#### Second Tile Load for Block (0,0)



#### Second Tile Use for Block (0,0), k of 0

$N_{0,0}$	N <sub>0,1</sub>	N <sub>0,2</sub>	
N <sub>1,0</sub>	N <sub>1,1</sub>	N <sub>1,2</sub>	
N <sub>2,0</sub>		N <sub>2,2</sub>	

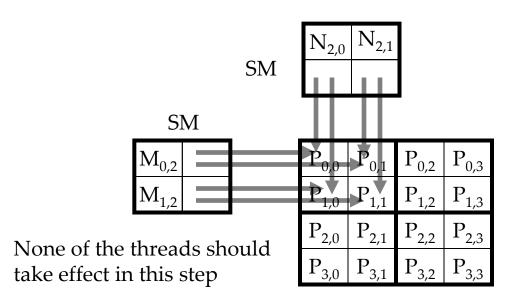
$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	
$M_{1,0}$	$M_{1,1}$	$M_{1,2}$	
$M_{2,0}$	$M_{2,1}$	$M_{2,2}$	



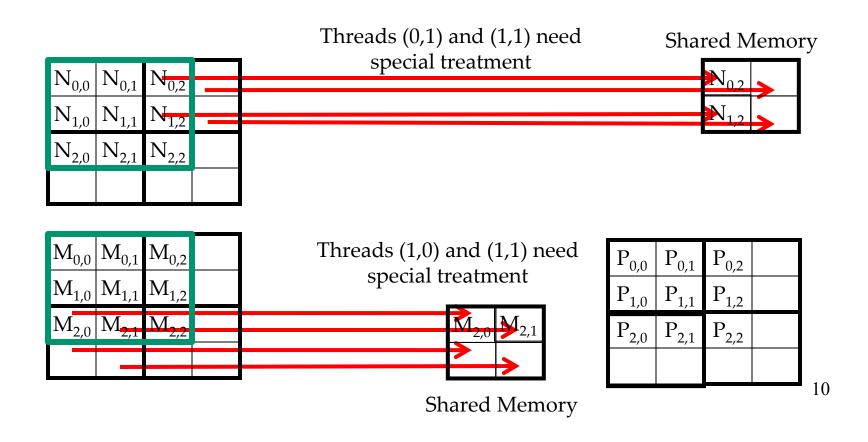
#### Second Tile Use for Block (0,0), k of 1

N <sub>0,0</sub>	N <sub>0,1</sub>	N <sub>0,2</sub>	

$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	
$M_{1,0}$	$M_{1,1}$	$M_{1,2}$	
$M_{2,0}$	$M_{2,1}$	$M_{2,2}$	

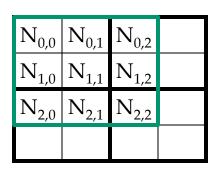


### First Tile Load for Block (1,1)

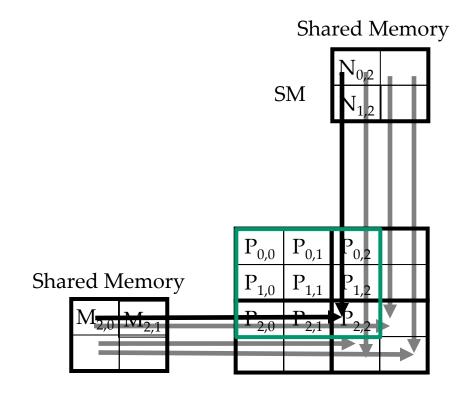


© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/ University of Illinois at Urbana-Champaign

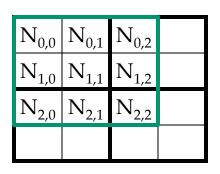
### First Tile Use for Block (1,1), k of 0



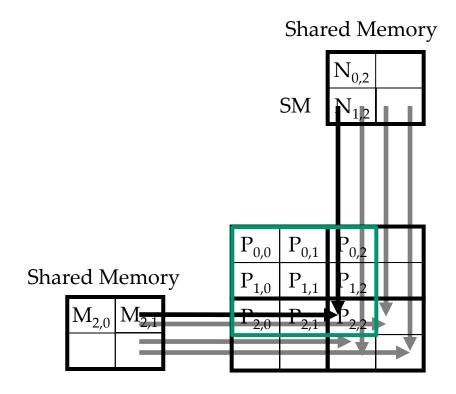
$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	
$M_{1,0}$	$M_{1,1}$	$M_{1,2}$	
$M_{2,0}$	$M_{2,1}$	$M_{2,2}$	



#### First Tile Use for Block (1,1), k of 1



$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	
$M_{1,0}$	$M_{1,1}$	$M_{1,2}$	
$M_{2,0}$	$M_{2,1}$	$M_{2,2}$	



#### Major Cases in Toy Example

- Threads that calculate valid P elements but can step outside valid input
  - Second tile of Block(0,0), all threads when k is 1
- Threads that do not calculate valid P elements
  - Block(1,1), Thread(1,0), non-existent row
  - Block(1,1), Thread(0,1), non-existing column
  - Block(1,1), Thread(1,1), non-existing row/column

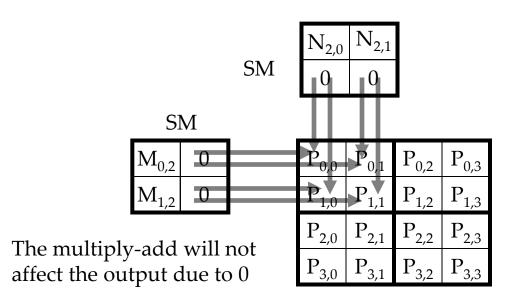
#### Solution: Write 0 for Missing Elements

- Test during tile load:
   is target within input matrix?
  - If yes, proceed to load;
  - otherwise, just write 0 to shared memory.
- The benefit?
  - No specialization during tile use!
  - Multiplying by 0 guarantees that unwanted terms do not contribute to the inner product.

### Second Tile Use for Block (0,0), k of 1

N <sub>0,0</sub>	N <sub>0,1</sub>	N <sub>0,2</sub>	

$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	
$M_{1,0}$	$M_{1,1}$	$M_{1,2}$	
$M_{2,0}$	$M_{2,1}$	$M_{2,2}$	

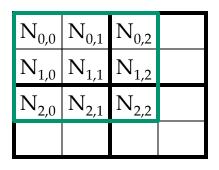


#### What About Threads Outside of P?

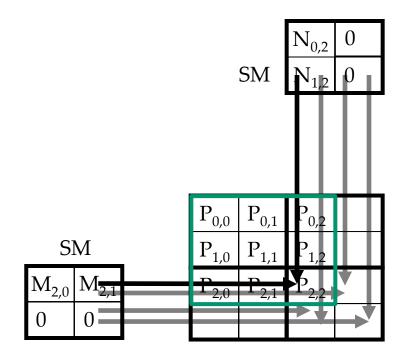
- If a thread is not within P,
  - All terms in sum are 0.
  - No harm in performing FLOPs.
  - No harm in writing to registers.
  - Must not be allowed to write to global memory!

So: Threads outside of P calculate 0, but store nothing.

#### First Tile Use for Block (1,1)



$M_{0,0}$	$M_{0,1}$	$M_{0,2}$	
$M_{1,0}$	$M_{1,1}$	$M_{1,2}$	
$M_{2,0}$	$M_{2,1}$	$M_{2,2}$	



#### Modifying the Tile Count

```
8. for (int m = 0; m < Width/TILE_WIDTH; ++m) {</pre>
```

The bound for m implicitly assumes that Width is a multiple of TILE\_WIDTH. We need to round up.

```
for (int m = 0; m < (Width - 1)/TILE_WIDTH + 1; ++m) {
```

For non-multiples of TILE\_WIDTH:

- quotient is unchanged;
- add one to round up.

For multiples of **TILE\_WIDTH**:

- quotient is now one smaller,
- but we add 1.

#### Modifying the Tile Loading Code

We had ...

```
// Collaborative loading of M and N tiles into shared memory
9. subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];
10. subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
```

Note: the tests for M and N tiles are NOT the same.

```
if (Row < Width && m*TILE_WIDTH+tx < Width) {
    // as before
    subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];
} else {
    subTileM[ty][tx] = 0;
}</pre>
```

#### And for Loading N...

#### We had ...

```
// Collaborative loading of M and N tiles into shared memory
9. subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];
10. subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
```

Note: the tests for M and N tiles are NOT the same.

```
if (m*TILE_WIDTH+ty < Width && Col < Width ) {
    // as before
    subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
} else {
    subTileN[ty][tx] = 0;
}</pre>
```

#### Modifying the Tile Use Code

We had ...

12. for (int k = 0; k < TILE\_WIDTH; ++k)

13. Pvalue += subTileM[ty][k] \* subTileN[k][tx];

Note: no changes are needed, but we might save a little energy (fewer floating-point ops)?

Pvalue += subTileM[ty][k] \* subTileN[k][tx];

(Row < Width && Col < Width) {

for (int k = 0; k < TILE WIDTH; ++k)

21

// as before

#### Modifying the Write to P

```
We had ...
16. P[Row*Width+Col] = Pvalue;
We must test for threads outside of P:
  (Row < Width && Col < Width) {
    // as before
    P[Row*Width+Col] = Pvalue;
```

#### Some Important Points

- For each thread, conditions are different for
  - Loading M element
  - Loading N element
  - Calculation/storing output elements
- Branch divergence
  - affects only blocks on boundaries, and
  - should be small for large matrices.
- What about rectangular matrices?

### Global Memory (DRAM) Bandwidth

Ideal



Reality



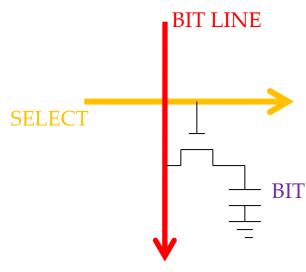
#### Most Large Memories Use DRAM

 Random Access Memory (RAM): same time needed to read/write any address

Dynamic RAM (DRAM):

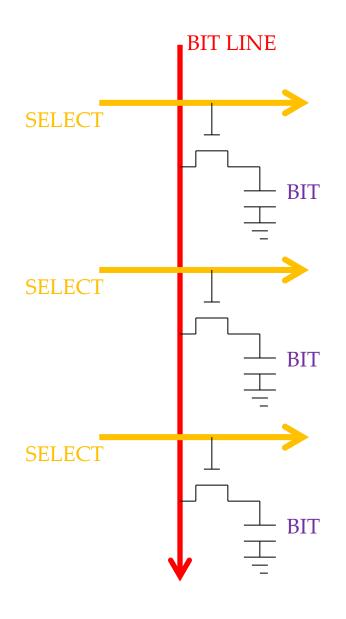
- bit stored on a capacitor
- connected via transistor to
   bit line for read/write
- bits disappear after a while

   (around 50 msec, due to tiny
   leakage currents through transistor),
   and must be rewritten (hence dynamic)



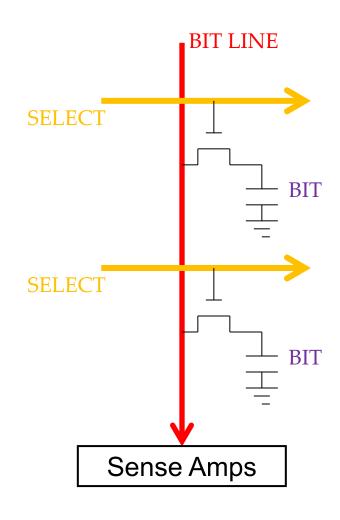
#### Many Cells (Bits) per Bit Line

- About 1,000 cells connect to each BIT LINE.
- Connection/disconnection depends on SELECT line.
- Some address bits decoded to connect exactly one cell to the BIT LINE.

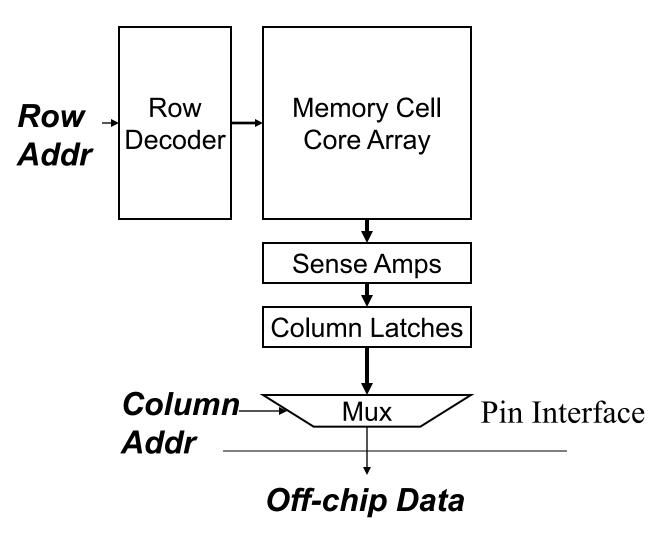


#### DRAM is Slow But Dense

- Capacitance...
  - tiny for the BIT, but
  - huge for theBIT LINE
- Use an amplifier for higher speed!
- Still slow...
- But only need
   1 transistor per bit.



#### **DRAM Bank Organization**

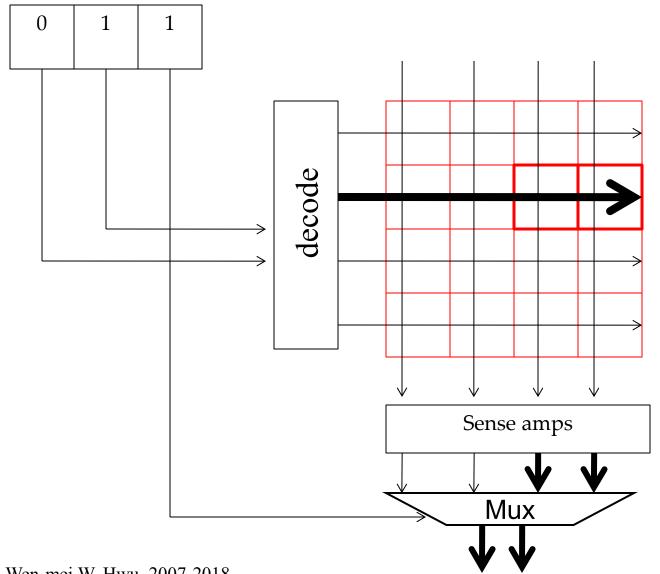


- SELECT lines connect to about 1,000 bit lines.
- Core array has about O(1M) bits
- Use more address bits to choose bit line(s).

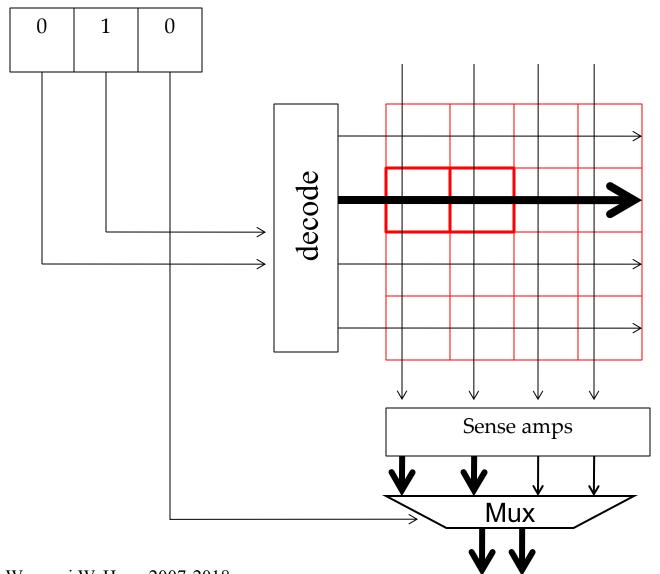
#### DRAM Interfaces are Clocked

- DRAM cells are not clocked (clocking requires transistors).
- DRAM interfaces are clocked.
  - DDR: Core speed = ½ interface speed
  - DDR2/GDDR3: Core speed = ¼ interface speed
  - DDR3/GDDR4: Core speed = ½ interface speed
  - likely to be worse in the future

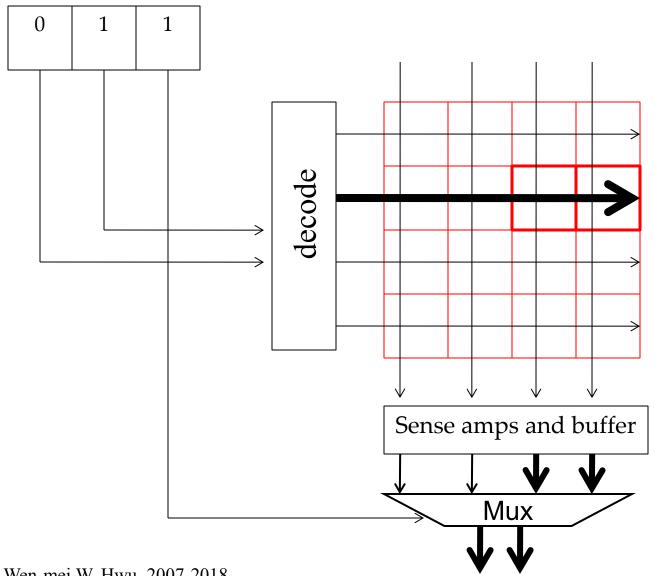
#### A very small (8x2 bit) DRAM Bank



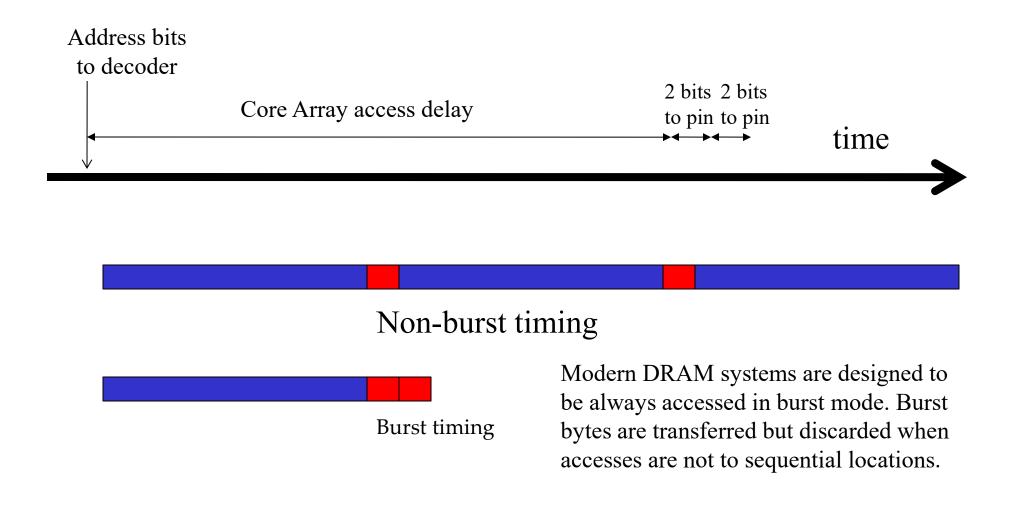
### DRAM Bursting (burst size = 4 bits)



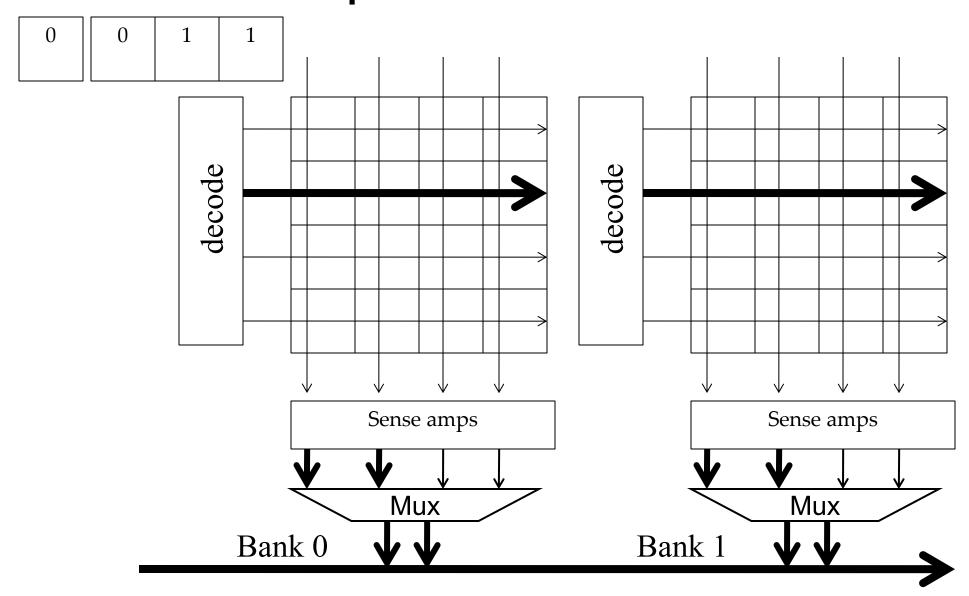
#### Second part of the burst



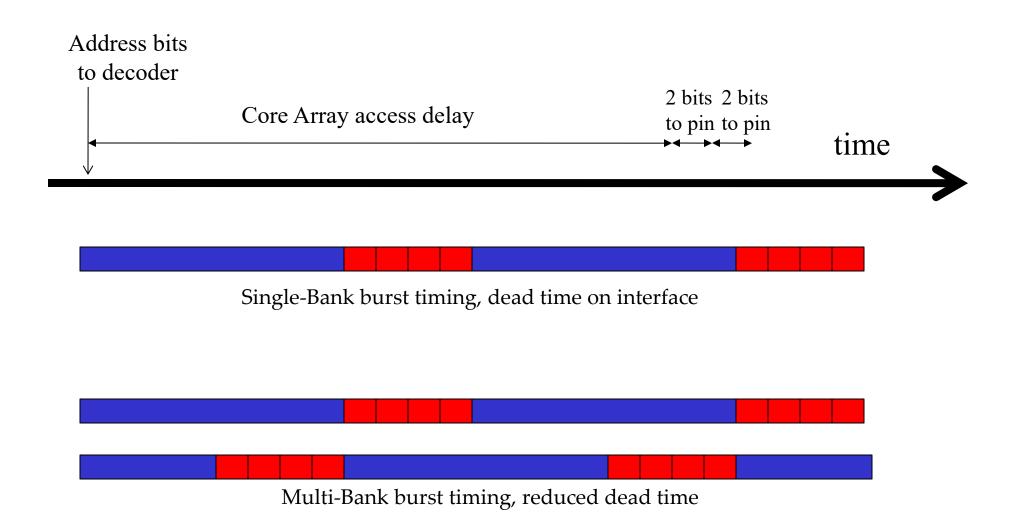
#### DRAM Bursting for the 8x2 Bank



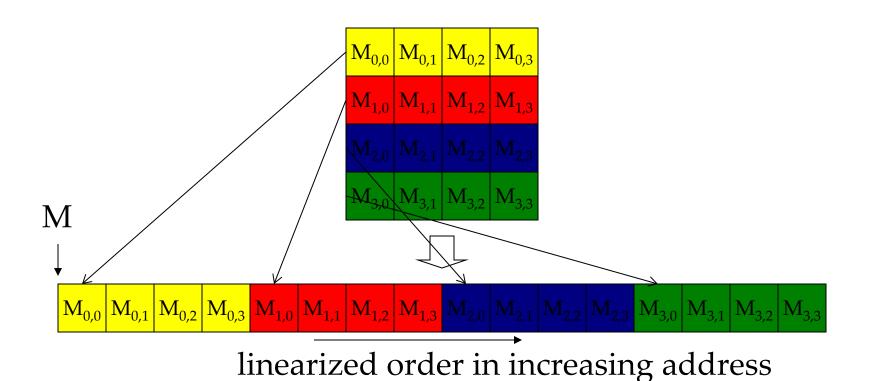
#### Multiple DRAM Banks



#### DRAM Bursting for the 8x2 Bank



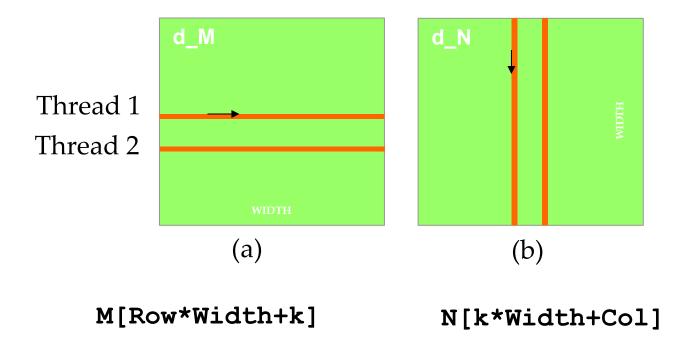
# Placing a 2D C array into linear memory space (review)



# A Simple Matrix Multiplication Kernel (review)

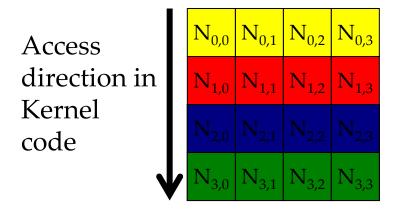
```
global void MatrixMulKernel(float* M, float* N, float* P, int Width)
// Calculate the row index of the P element and M
int Row = blockIdx.y * blockDim.y + threadIdx.y;
// Calculate the column index of P and N
int Col = blockIdx.x * blockDim.x + threadIdx.x;
if ((Row < Width) && (Col < Width)) {
   float Pvalue = 0;
   // each thread computes one element of the block sub-matrix
   for (int k = 0; k < Width; ++k)
     Pvalue += M[Row*Width+k] * N[k*Width+Col];
   P[Row*Width+Col] = Pvalue;
```

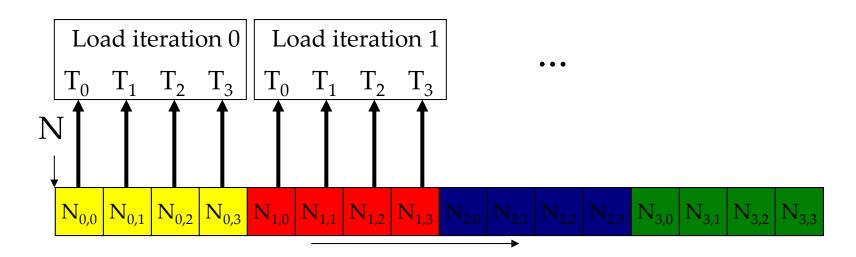
#### Two Access Patterns



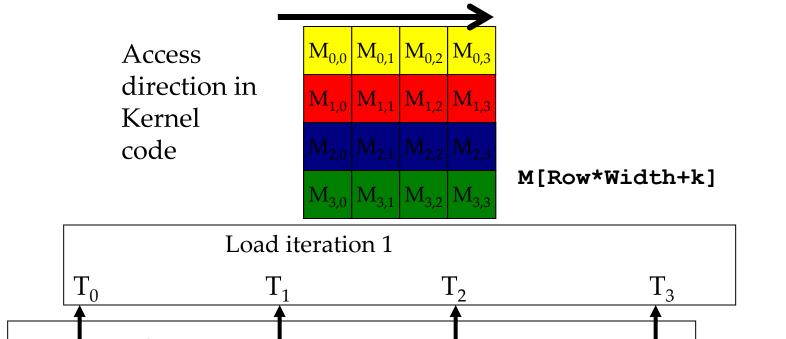
k is loop counter in the inner product loop of the kernel code

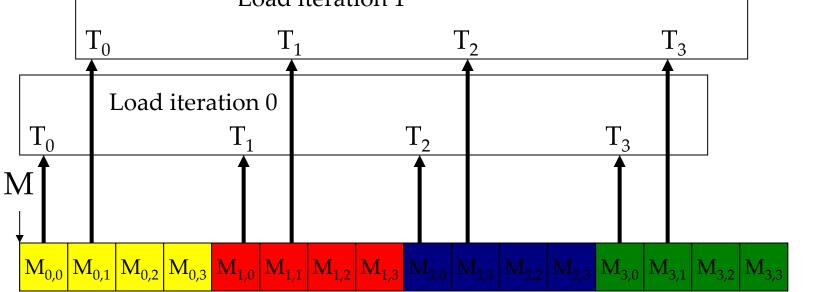
#### N accesses are coalesced





#### M accesses are not coalesced



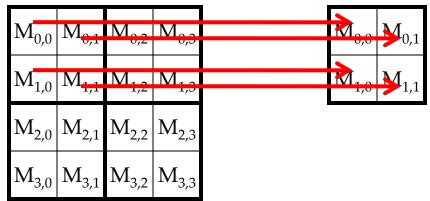


#### Work for Block (0,0) Step 0

#### Shared Memory

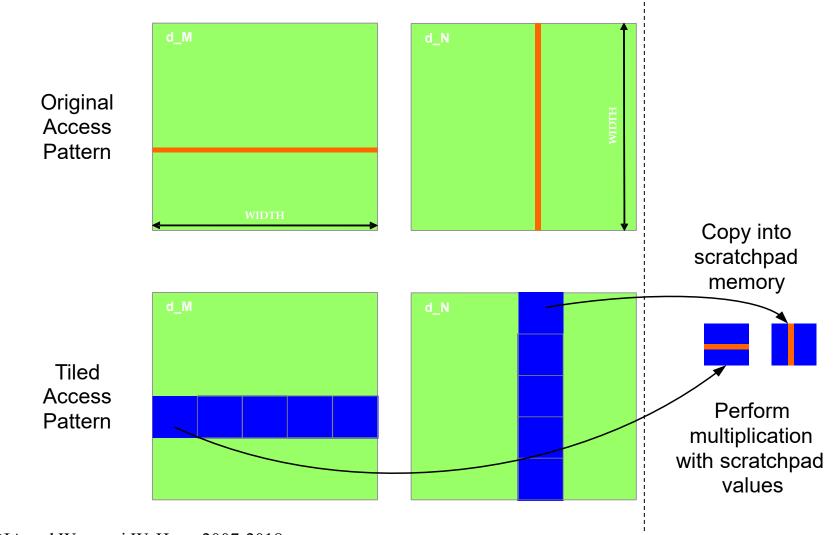


#### Shared Memory



P <sub>0,0</sub>	P <sub>0,1</sub>	P <sub>0,2</sub>	P <sub>0,3</sub>
P <sub>1,0</sub>	P <sub>1,1</sub>	P <sub>1,2</sub>	P <sub>1,3</sub>
P <sub>2,0</sub>	P <sub>2,1</sub>	P <sub>2,2</sub>	P <sub>2,3</sub>
P <sub>3,0</sub>	P <sub>3,1</sub>	P <sub>3,2</sub>	P <sub>3,3</sub>

# Use shared memory to enable coalescing in tiled matrix multiplication



#### ANY MORE QUESTIONS? READ CHAPTER 5

#### **Problem Solving**

 Q: Which of the following two lines of a kernel code has better performance?

```
int tx = blockIdx.x * blockDim.x + threadIdx.x;
int ty = blockIdx.y * blockDim.y + threadIdx.y;
A) B[ty * Width + tx] = 2 * A[ty * Width + tx];
B) B[tx * Width + ty] = 2 * A[tx * Width + ty];
```

A: Cannot tell. We do not know how the kernel was launched.

#### **Problem Solving**

• Q: Consider a DRAM system with a burst size of 512 bytes and a peak bandwidth of 240 GB/s. Assume a thread block size of 1024 and warp size of 32 and that A is a floating-point array in the global memory. What is the maximal memory data access throughput we can hope to achieve in the following access to A?

```
int i = blockIdx.x * blockDim.x + threadIdx.x;
float temp = A[4*i] + A[4*i+1];
```

- A:
  - From a burst of 512 bytes, we will only use every other 8 bytes due to reading from [4\*i] and [4\*i+1] memory locations.
  - So, we can expect 120 GB/s.