



ECE408/CS483/CSE408 Spring 2025

## Applied Parallel Programming

# Lecture 16

## Atomic Operations and Histogramming

# Course Reminders

- Lab 6 is due this Friday
- Project Milestone 2 is out and is due in  $< 3$  weeks
- Take a note of the date/time for second Midterm Exam
  - May 6<sup>th</sup>

# Objective

- To understand atomic operations
  - Read-modify-write in parallel computation
  - A primitive form of “critical regions” in parallel programs
  - Use of atomic operations in CUDA
  - Why atomic operations reduce memory system throughput
  - How to avoid atomic operations in some parallel algorithms
- To learn practical histogram programming techniques
  - Basic histogram algorithm using atomic operations
  - Atomic operation throughput
  - Privatization

# A Common Collaboration Pattern

- Multiple bank tellers count the total amount of cash in the safe
- Each teller grabs a pile and counts
- Have a central display of the running total
- Whenever someone finishes counting a pile, add the subtotal of the pile to the running total
- A bad outcome
  - Some of the piles were not accounted for
  - How/why is this possible?

# A Common Arbitration Pattern

- Multiple customers booking air tickets
- Each
  - Brings up a flight seat map
  - Decides on a seat
  - Update the seat map, mark the seat as taken
- A bad outcome
  - Multiple passengers ended up booking the same seat
  - How/why is this possible?

# Read-Modify-Write Operations

thread1:  $\text{Old} \leftarrow \text{Mem}[x]$   
 $\text{New} \leftarrow \text{Old} + 1$   
 $\text{Mem}[x] \leftarrow \text{New}$

thread2:  $\text{Old} \leftarrow \text{Mem}[x]$   
 $\text{New} \leftarrow \text{Old} + 1$   
 $\text{Mem}[x] \leftarrow \text{New}$

If  $\text{Mem}[x]$  was **initially 0**, what would the value of  $\text{Mem}[x]$  be after threads 1 and 2 have completed?

– What does each thread get in their Old variable?

The answer may vary due to data races. To avoid data races, you should use atomic operations.

# Timing Scenario #1

Time	Thread 1	Thread 2
1	(0) Old $\leftarrow$ Mem[x]	
2	(1) New $\leftarrow$ Old + 1	
3	(1) Mem[x] $\leftarrow$ New	
4		(1) Old $\leftarrow$ Mem[x]
5		(2) New $\leftarrow$ Old + 1
6		(2) Mem[x] $\leftarrow$ New

- Thread 1 Old = 0
- Thread 2 Old = 1
- Mem[x] = 2 after the sequence

# Timing Scenario #2

Time	Thread 1	Thread 2
1		(0) Old $\leftarrow$ Mem[x]
2		(1) New $\leftarrow$ Old + 1
3		(1) Mem[x] $\leftarrow$ New
4	(1) Old $\leftarrow$ Mem[x]	
5	(2) New $\leftarrow$ Old + 1	
6	(2) Mem[x] $\leftarrow$ New	

- Thread 1 Old = 1
- Thread 2 Old = 0
- Mem[x] = 2 after the sequence



# Timing Scenario #3

Time	Thread 1	Thread 2
1	(0) Old $\leftarrow$ Mem[x]	
2	(1) New $\leftarrow$ Old + 1	
3		(0) Old $\leftarrow$ Mem[x]
4	(1) Mem[x] $\leftarrow$ New	
5		(1) New $\leftarrow$ Old + 1
6		(1) Mem[x] $\leftarrow$ New

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

# Timing Scenario #4

Time	Thread 1	Thread 2
1		(0) Old $\leftarrow$ Mem[x]
2		(1) New $\leftarrow$ Old + 1
3	(0) Old $\leftarrow$ Mem[x]	
4		(1) Mem[x] $\leftarrow$ New
5	(1) New $\leftarrow$ Old + 1	
6	(1) Mem[x] $\leftarrow$ New	

- Thread 1 Old = 0
- Thread 2 Old = 0
- Mem[x] = 1 after the sequence

# Atomic Operations Prevent Interleaving

thread1:  $\text{Old} \leftarrow \text{Mem}[x]$   
 $\text{New} \leftarrow \text{Old} + 1$   
 $\text{Mem}[x] \leftarrow \text{New}$

thread2:  $\text{Old} \leftarrow \text{Mem}[x]$   
 $\text{New} \leftarrow \text{Old} + 1$   
 $\text{Mem}[x] \leftarrow \text{New}$

Or

thread1:  $\text{Old} \leftarrow \text{Mem}[x]$   
 $\text{New} \leftarrow \text{Old} + 1$   
 $\text{Mem}[x] \leftarrow \text{New}$

thread2:  $\text{Old} \leftarrow \text{Mem}[x]$   
 $\text{New} \leftarrow \text{Old} + 1$   
 $\text{Mem}[x] \leftarrow \text{New}$

# Without Atomic Operations

Mem[x] initialized to 0

thread1: Old  $\leftarrow$  Mem[x]

New  $\leftarrow$  Old + 1

Mem[x]  $\leftarrow$  New

thread2: Old  $\leftarrow$  Mem[x]

New  $\leftarrow$  Old + 1

Mem[x]  $\leftarrow$  New

- Both threads receive 0
- Mem[x] becomes 1

# Needed When Threads Write to Same Location

When **two threads**

- **may write to** the **same memory** location,
- the program may **need atomic operations**.

Sharing is **not always easy to recognize**...

- Do two insertions into a hash table share data?
- What about two graph node updates based on all of the nodes' neighbors?
- What if nodes are on same side of bipartite graph?

# What Exactly is “Atomic?”

To a high-energy photon, atoms are not.

Atomicity is ALWAYS with respect to something.

Two sections of code

- that execute atomically with respect to one another
- appear to the software as though
- the programs’ execution did not interleave at all.

# What Can Go Wrong?

Common failure mode:

- **Programmer *thinks* operations are independent.**
- Hasn't considered input data for which they are not.
- Or another programmer reuses code without understanding assumptions that imply independence.

Also: **atomicity does not constrain relative order.**

# Implementing Atomic Operations

- Many ISAs offer synchronization primitives,
  - **instructions with** one (or more) **address operands**
  - **that execute atomically with respect to one another** when used on the same address.
- Mostly read, modify, write operations
  - Bit test and set
  - Compare and swap / exchange
  - Swap / exchange
  - Fetch and increment / add



# Atomicity Enforced by Microarchitecture

When synchronization primitives execute,

- **hardware ensures** that **no other** thread
- **accesses** the location **until** the operation is **complete**.

Other threads that access the location

- are typically stalled or held in a queue until their turn.
- **Threads perform atomic operations serially.**

# Atomic Operations in CUDA

- Function calls that are translated into single ISA instructions (a.k.a. *intrinsics*)
  - Atomic add, sub, inc, dec, min, max, exch (exchange), CAS (compare and swap)
  - Read CUDA C programming Guide for more details

- Atomic Add

*int atomicAdd(int\* **address**, int **val**);*

reads the 32-bit word **old** pointed to by **address** in global or shared memory, computes (**old** + **val**), and stores the result back to memory at the same address. The function returns **old**.

***old** = \***address**; \***address** = **old** + **val**; return **old**;*

# More Atomic Adds in CUDA

- Unsigned 32-bit integer atomic add

*unsigned int atomicAdd(unsigned int\* address, unsigned int val);*

- Unsigned 64-bit integer atomic add

*unsigned long long int atomicAdd(unsigned long long int\* address, unsigned long long int val);*

- Single-precision floating-point atomic add  
(capability > 2.0)

– *float atomicAdd(float\* address, float val);*

# Building synchronization with atomics

- How would we build `__syncthreads()` for block?
- How would we create `__syncthreads()` for entire grid?
  - And why would this not be a good idea?
- How would we create a critical section? I.e., one thread per block executing a particular section of code?
- How would we create a critical section per grid?
  - Why doesn't this have the same issue as `__syncthreads()` for grid?

# Atomic Compare and Swap (CAS)

```
bool atomicCAS(int *address, int old, int new)
{
    if (*address != old)
        return false;

    *address = new;

    return true;
}
```

- CAS is an atomic instruction used in multithreading to achieve synchronization.
- It compares the contents of a memory location with a given value and, only if they are the same, modifies the contents of that memory location to a new given value. This is done as a single atomic operation.

# Atomic Add using CAS

```
int atomicAdd(int* address, int value)
{
    bool done = false;

    while (!done) {
        old_v = *address;
        done = atomicCAS(address, old_v, old_v+value);
    }

    return old_v;
}
```

- The function performs the action  $*\text{address} \leftarrow * \text{address} + \text{value}$  *atomically* and returns the original value stored at address.
- There is no requirement that any sequence of operations is atomic except for atomicCAS.

# Histogramming

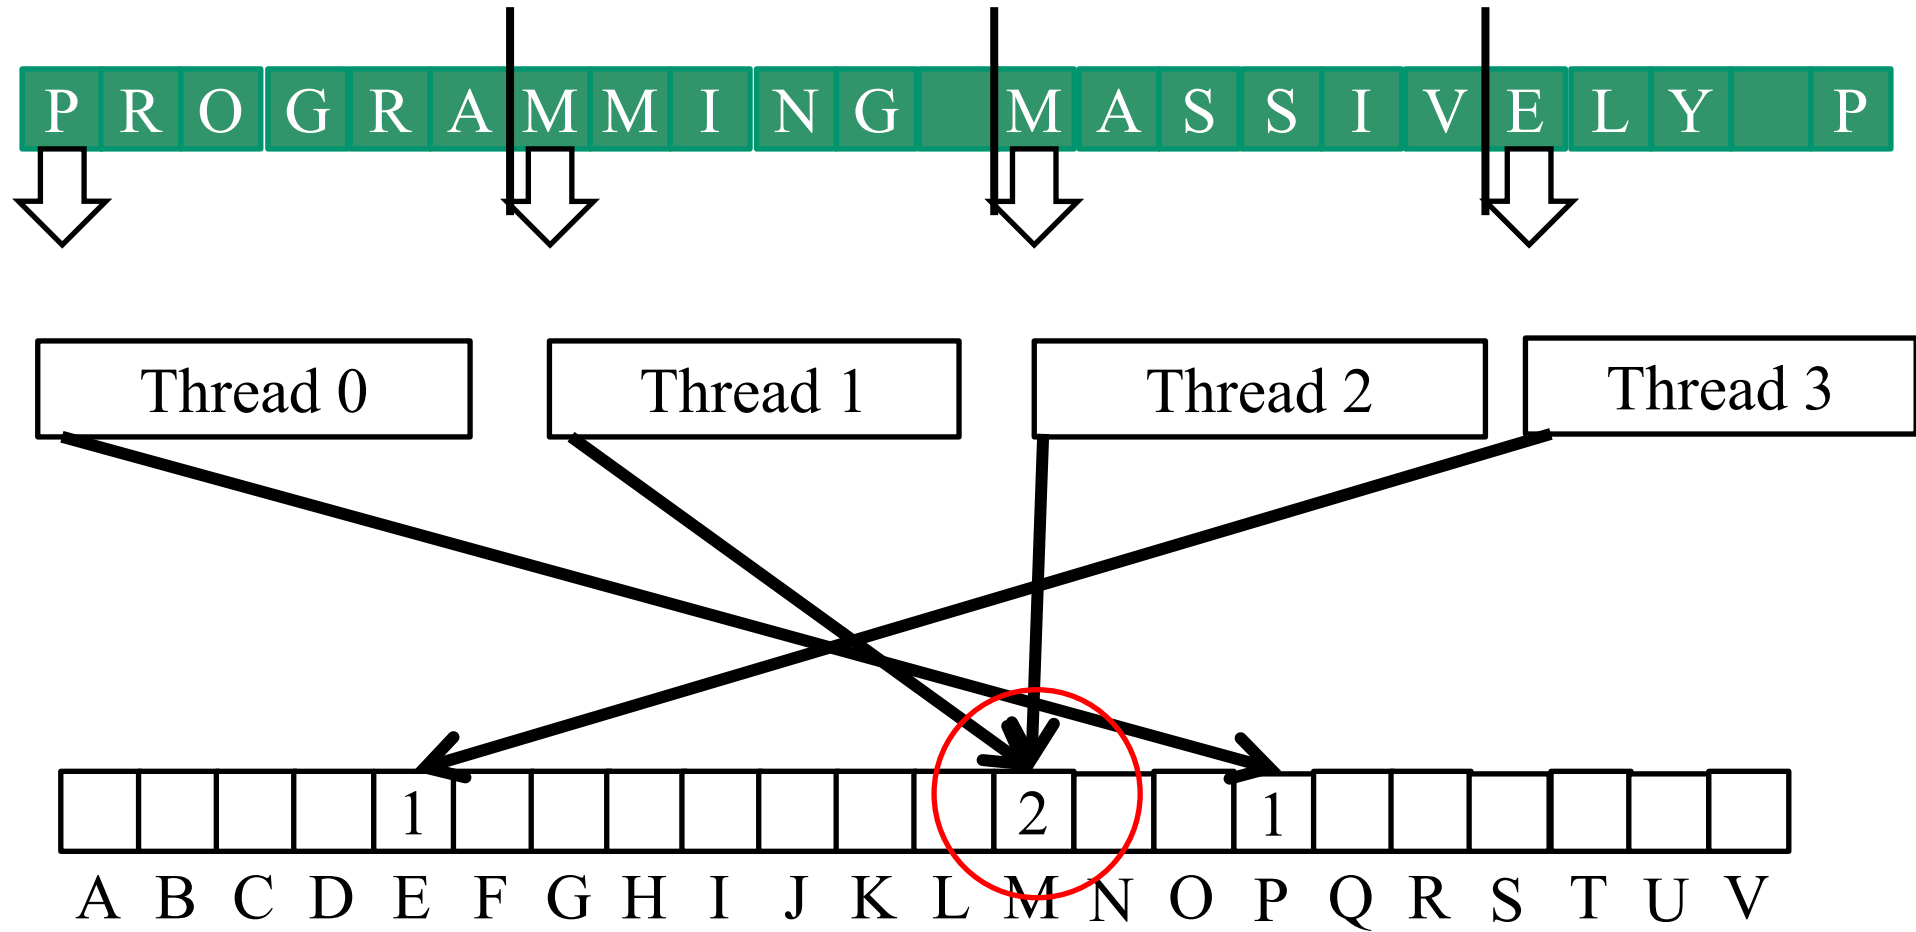
- A method for extracting notable features and patterns from large data sets
  - Feature extraction for object recognition in images
  - Fraud detection in credit card transactions
  - Correlating heavenly object movements in astrophysics
  - ...
- Basic histograms - for each element in the data set, use the value to identify a “bin” to increment

# A Histogram Example

- In sentence “Programming Massively Parallel Processors” build a histogram of frequencies of each letter
- A(4), C(1), E(1), G(1), ...
- How do you do this in parallel?

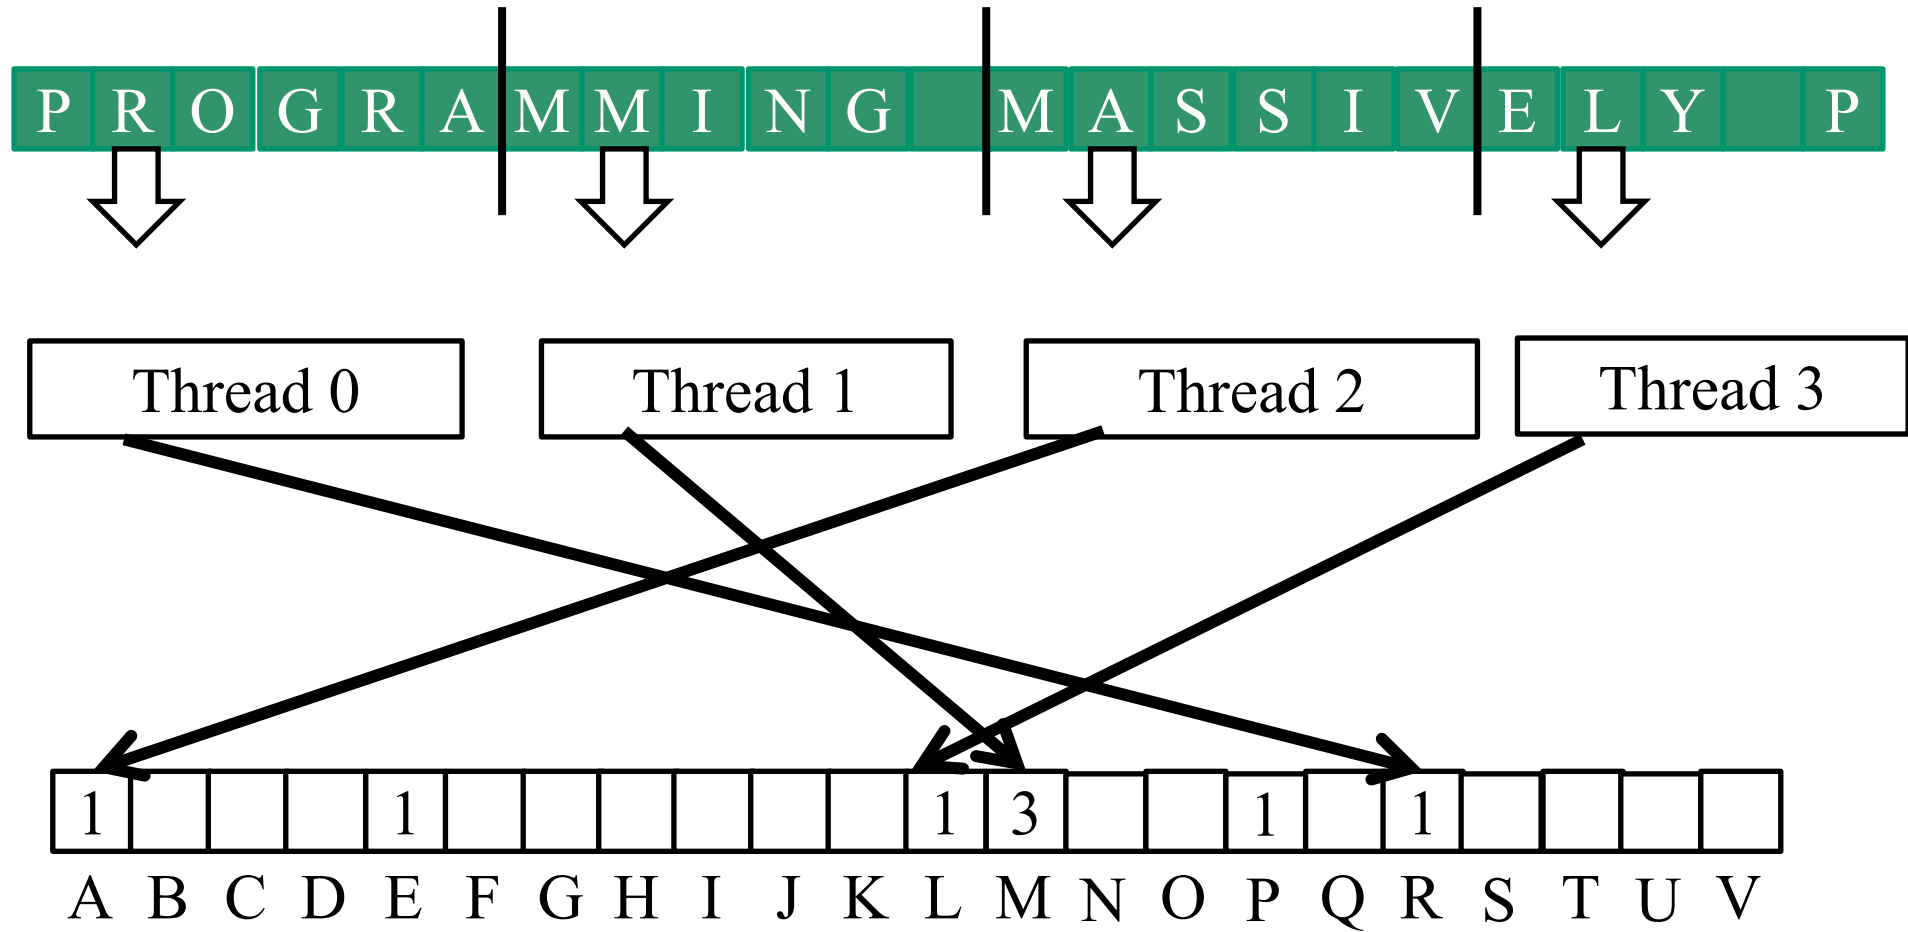


# Iteration #1 – 1<sup>st</sup> letter in each section

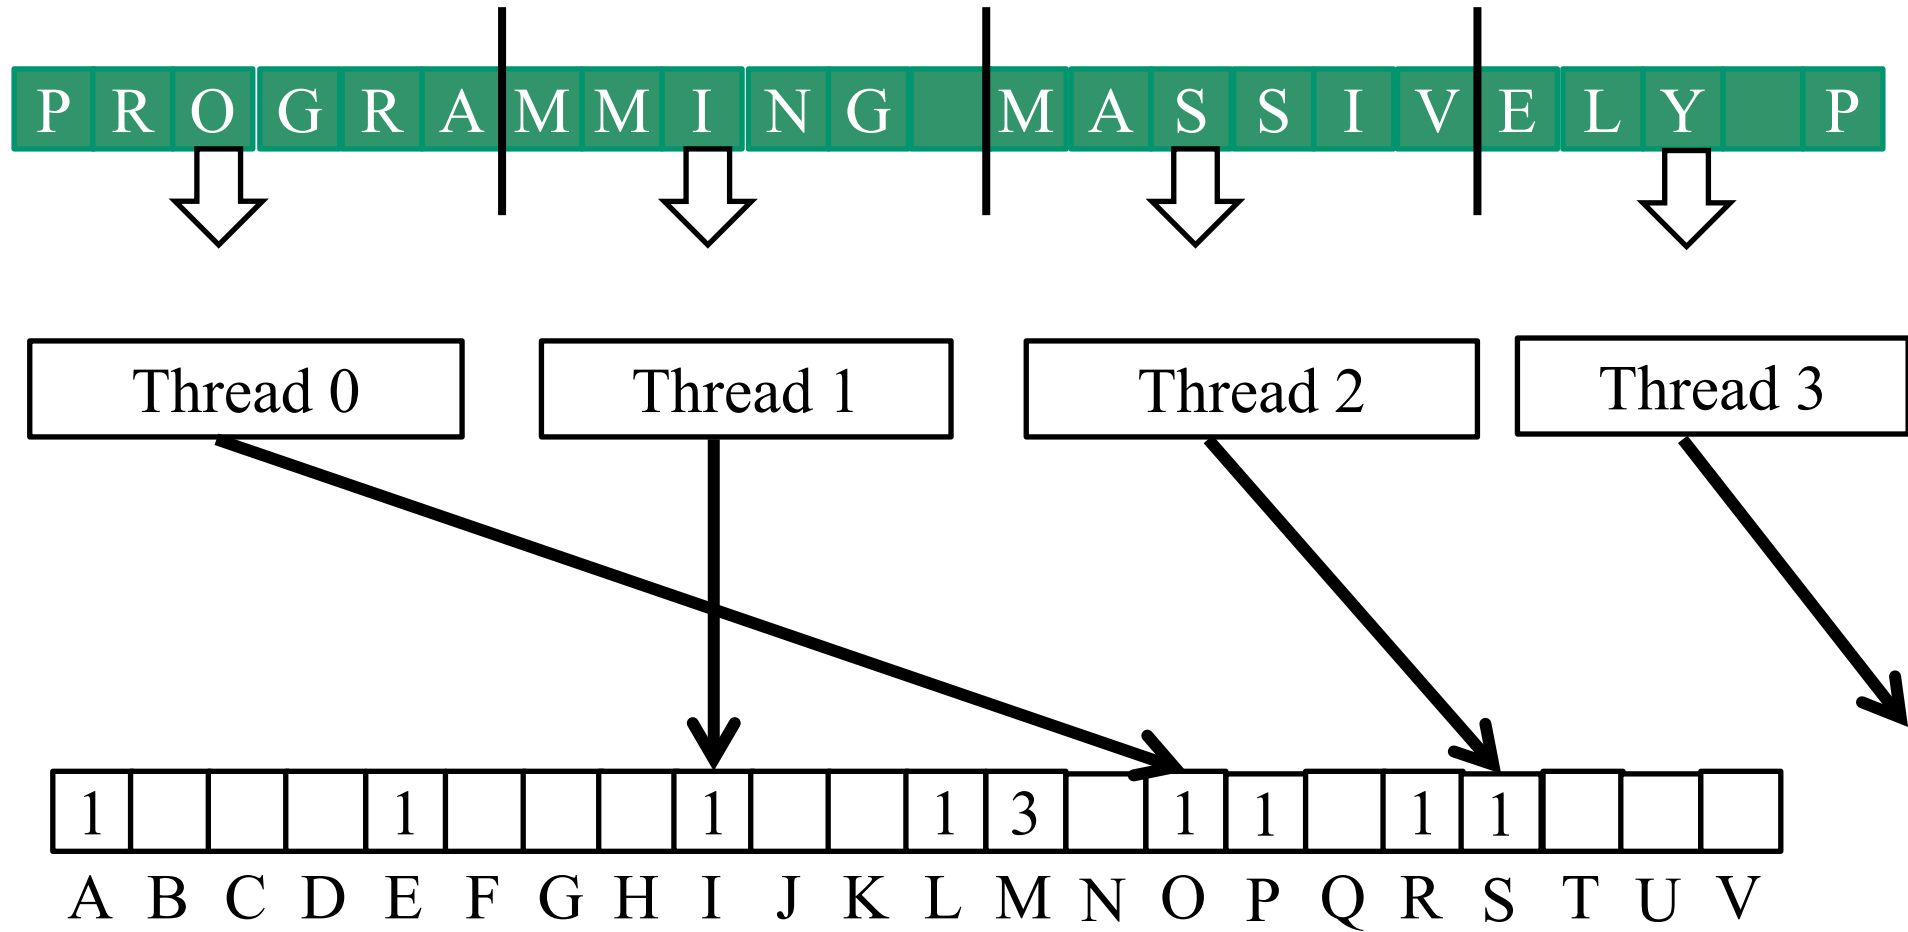


Atomic operation enforces correct update

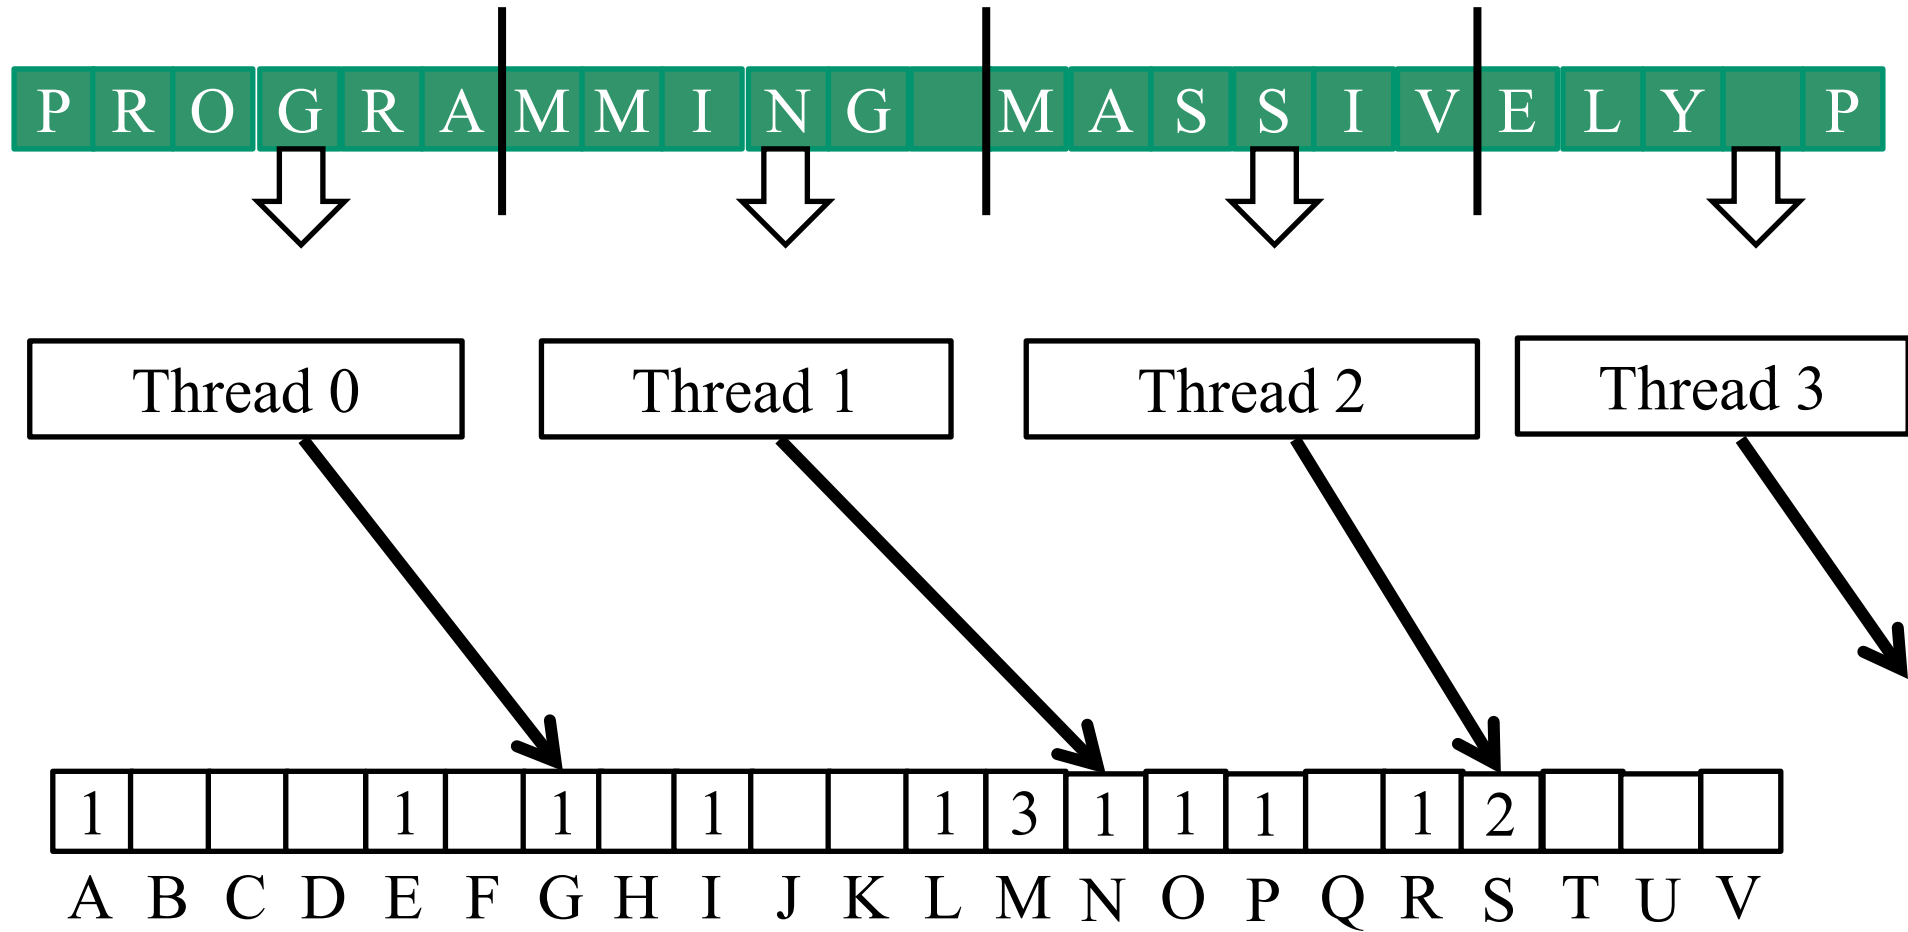
# Iteration #2 – 2<sup>nd</sup> letter in each section



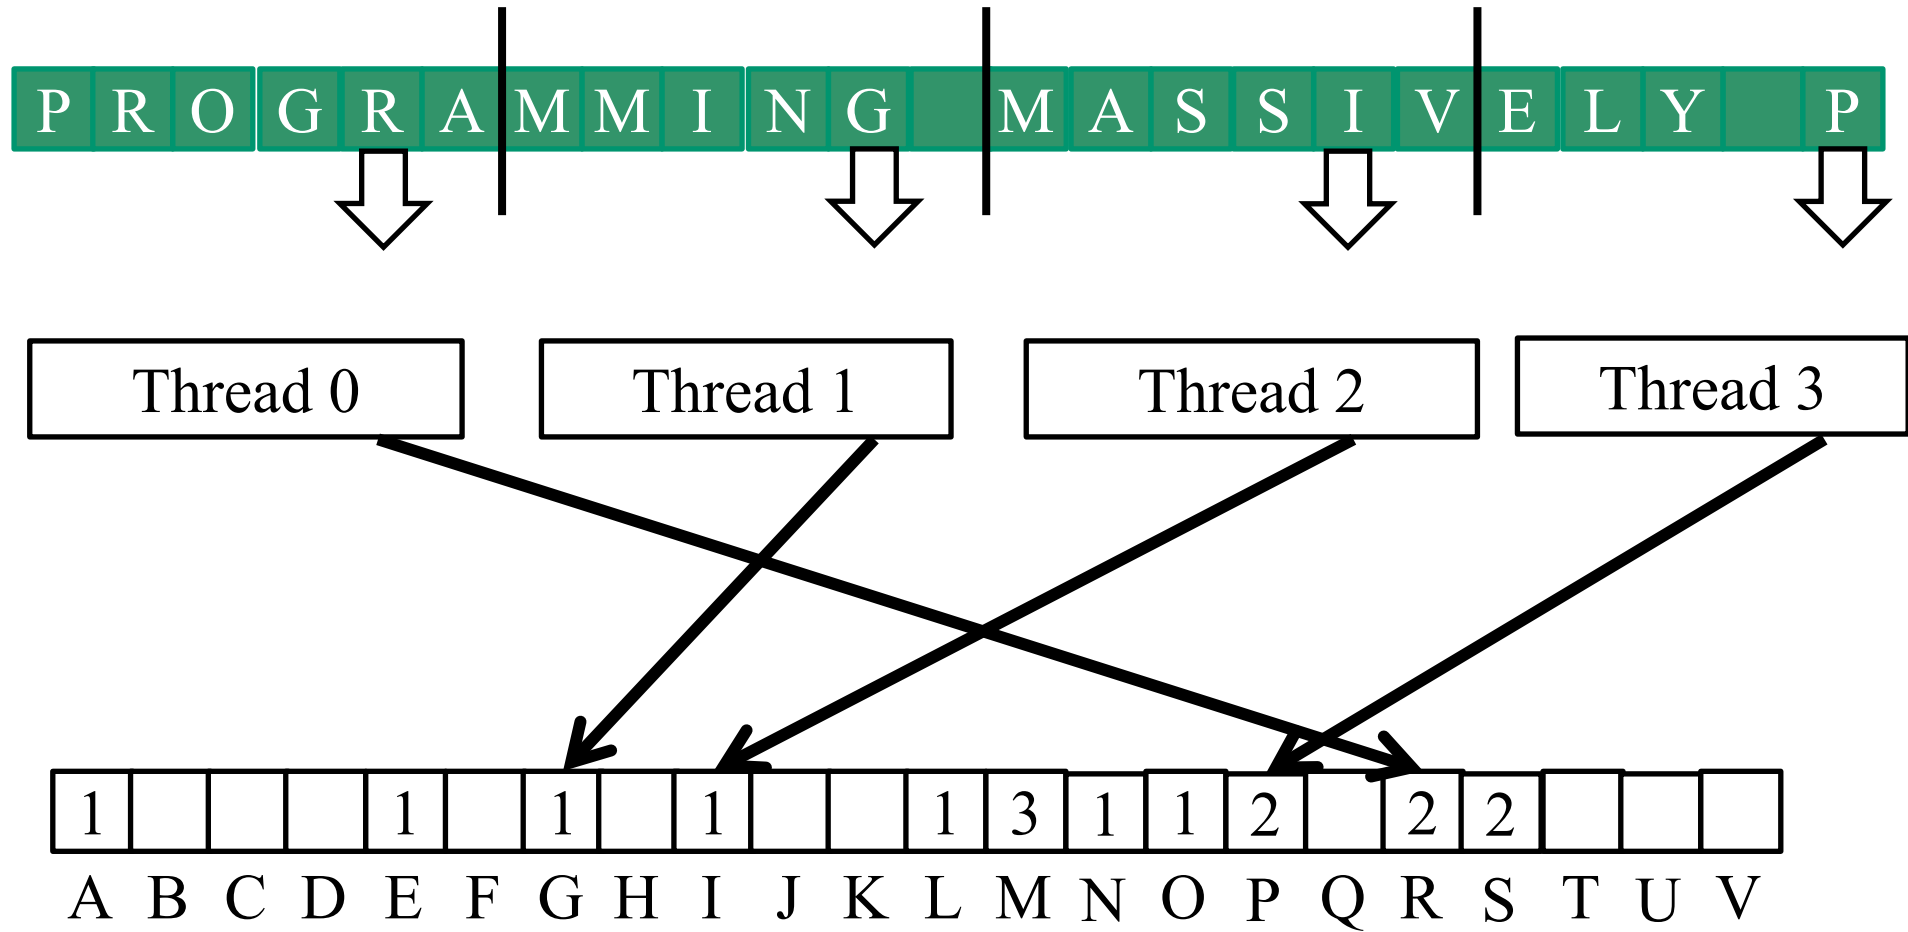
# Iteration #3



# Iteration #4

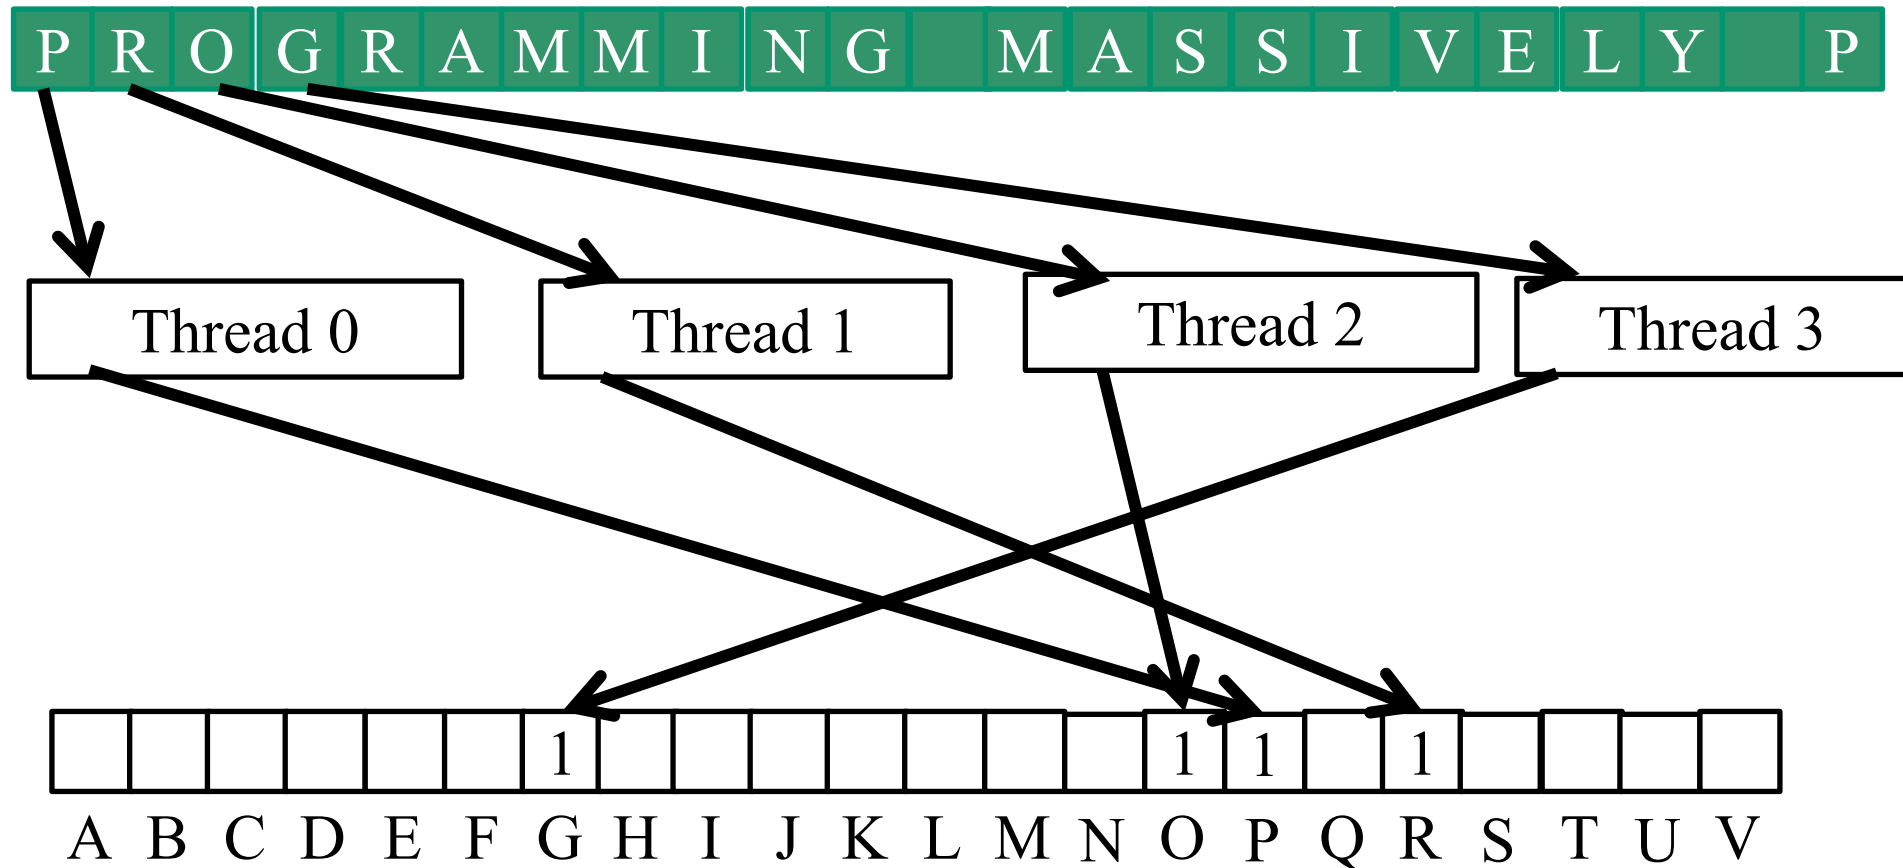


# Iteration #5



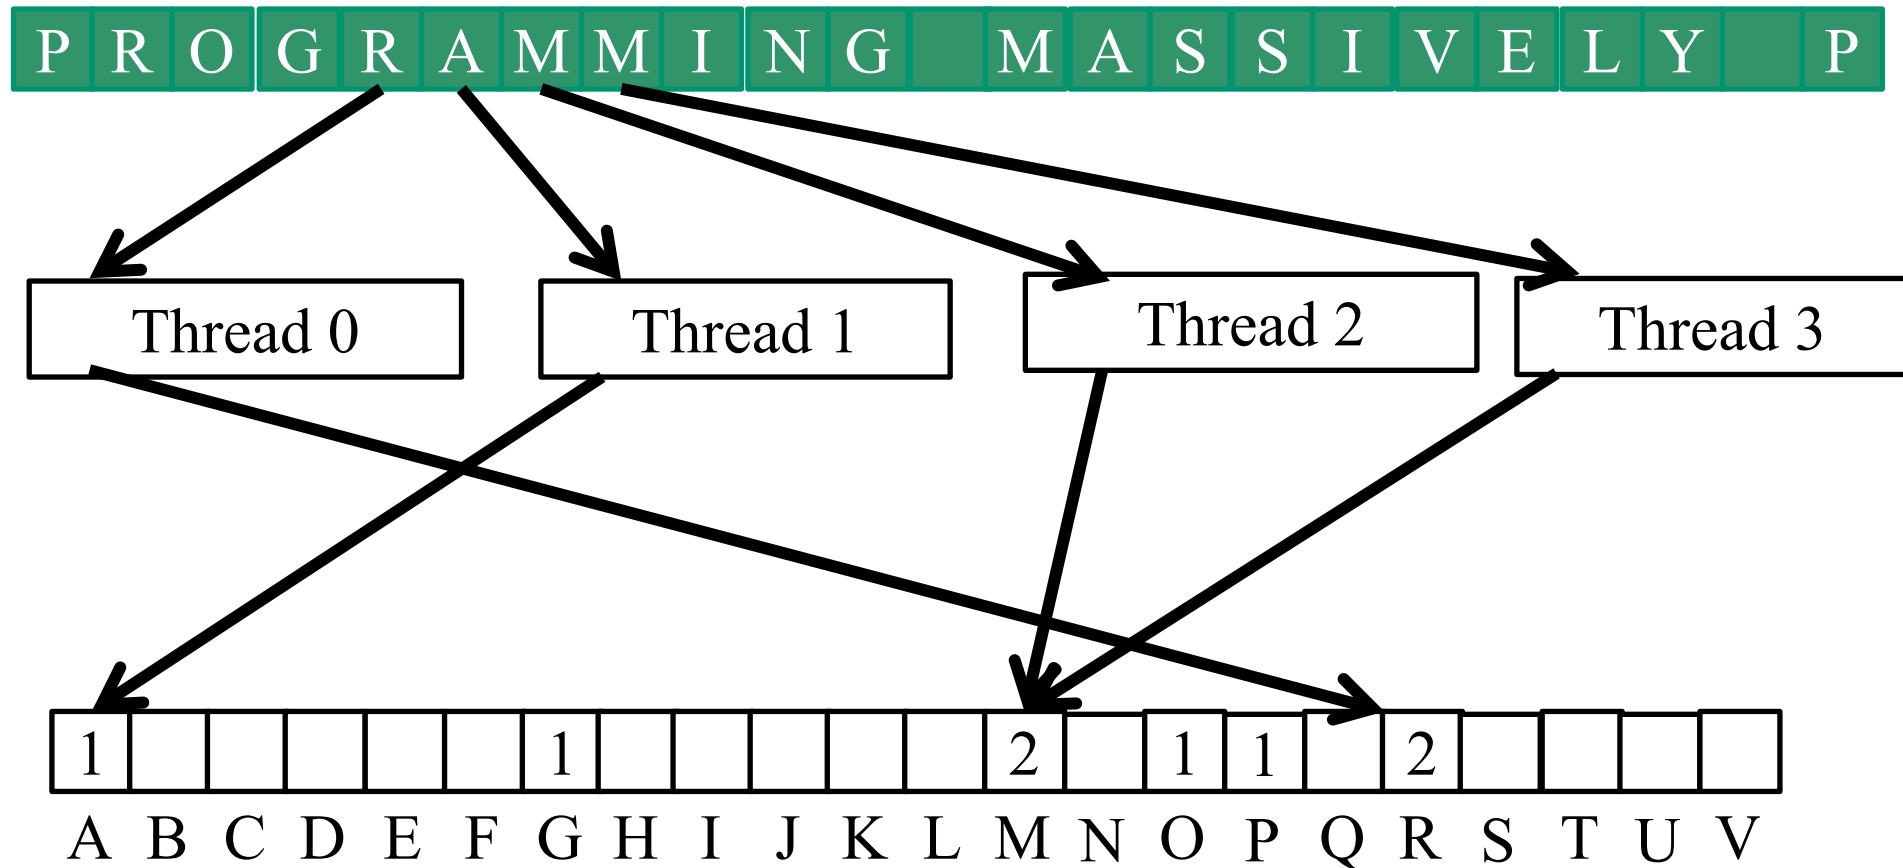
# A better approach

- Reads from the input array are not coalesced
  - Assign inputs to each thread in a strided pattern
  - Adjacent threads process adjacent input letters



# Iteration 2

- All threads move to the next section of input



# A Histogram Kernel

- The kernel receives a pointer to the input buffer
- Each thread process the input in a strided pattern

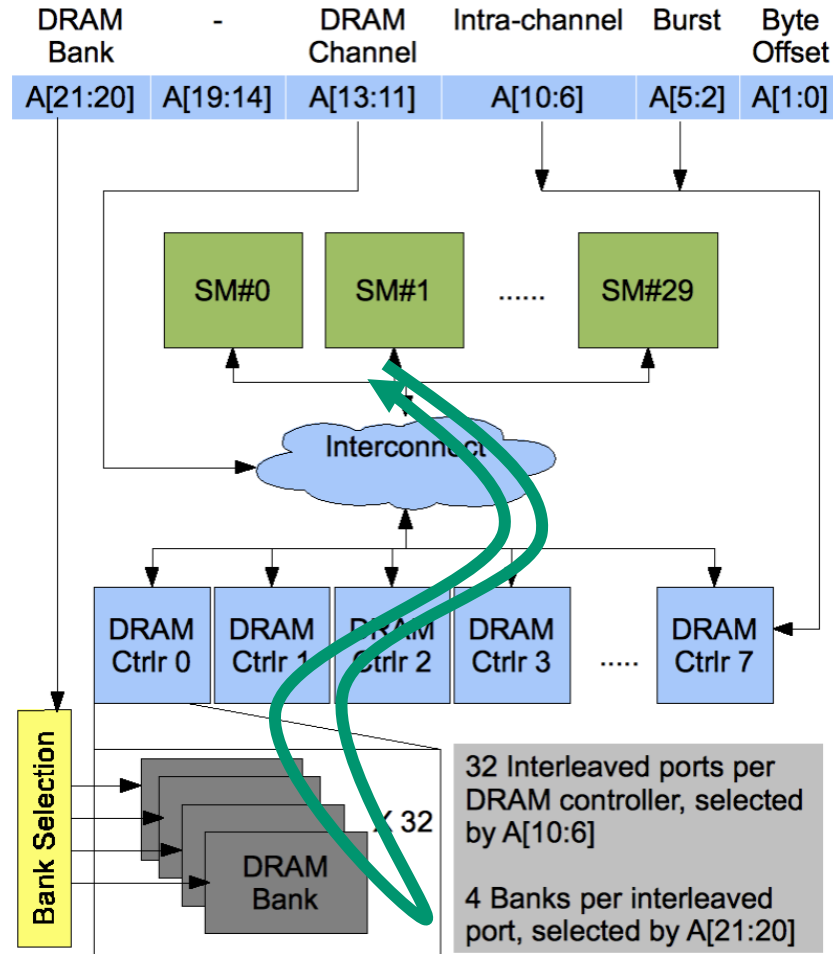
```
__global__ void histo_kernel(char *buf, long size, int *histo)
{
    int i = threadIdx.x + blockIdx.x * blockDim.x;
    int stride = blockDim.x * gridDim.x; // stride = total # of threads

    // All threads in the grid collectively handle
    // blockDim.x * gridDim.x consecutive elements

    while (i < size) {
        atomicAdd( &(histo[buf[i]]), 1);
        i += stride;
    }
}
```

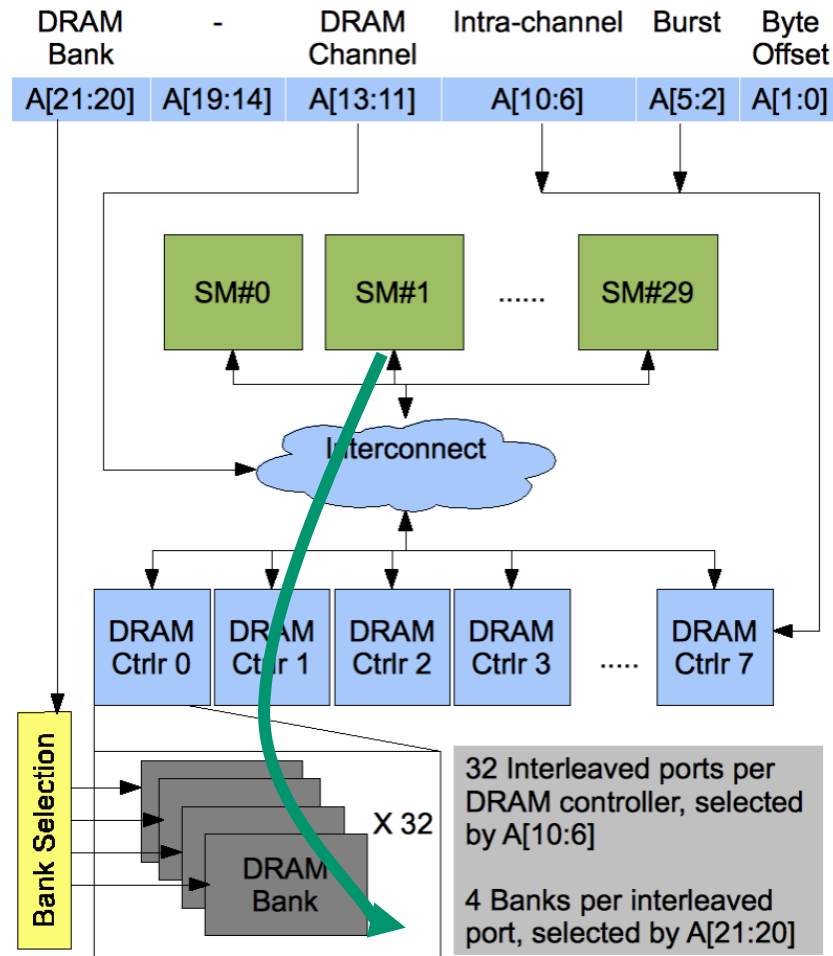


# Atomic Operations on DRAM



- An atomic operation starts with a read, with a latency of a few hundred cycles

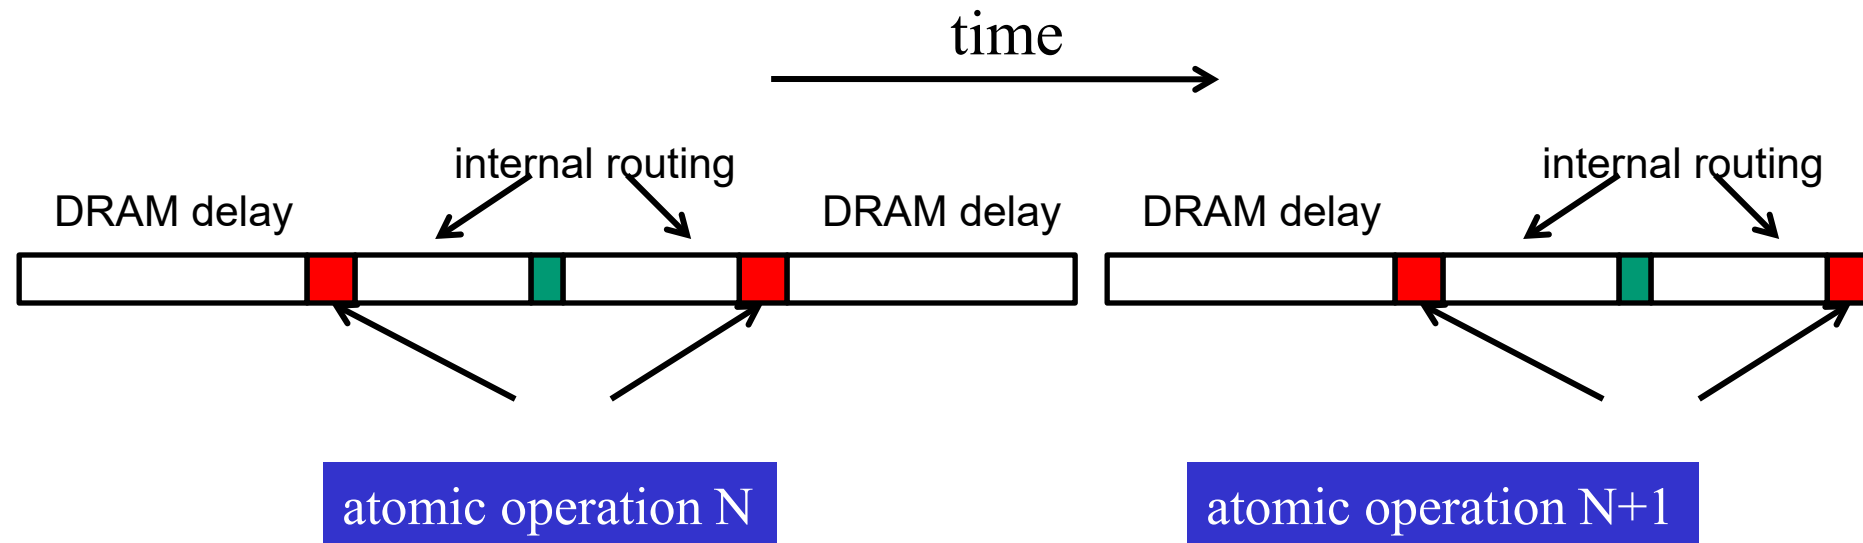
# Atomic Operations on DRAM



- An atomic operation starts with a read, with a latency of a few hundred cycles
- The atomic operation ends with a write, with a latency of a few hundred cycles
- During this whole time, no one else can access the location

# Atomic Operations on DRAM

- Each Load-Modify-Store has two full memory access delays
  - All atomic operations on the same variable (RAM location) are serialized



# Latency determines throughput of atomic operations


- Throughput of an atomic operation is the rate at which the application can execute an atomic operation on a particular location.
- The rate is limited by the total latency of the read-modify-write sequence, typically more than 1000 cycles for global memory (DRAM) locations.
- This means that if many threads attempt to do atomic operation on the same location (contention), the memory bandwidth is reduced to  $< 1/1000!$

# You may have a similar experience in supermarket checkout

- Some customers realize that they missed an item after they started to check out
- They run to the isle and get the item while the line waits
  - The rate of check is reduced due to the long latency of running to the isle and back.
- Imagine a store where every customer starts the check out before they even fetch any of the items
  - The rate of the checkout will be  $1 / (\text{entire shopping time of each customer})$


# Hardware Improvements

- Atomic operations on L2 cache
  - medium latency, but still serialized
  - Global to all blocks
  - “Free improvement” on Global Memory atomics



# Hardware Improvements

- Atomic operations on Shared Memory
  - Very short latency, but still serialized
  - Private to each thread block
  - Need algorithm work by programmers (more later)



# Atoms in Shared Memory Requires Privatization

- Create private copies of the histo[] array for each thread block

```
__global__ void histo_kernel(unsigned char *buffer,
                             long size, unsigned int *histo)
{
    __shared__ unsigned int histo_private[256];
    // warning: this will not work correctly if there are fewer than 256 threads!
    if (threadIdx.x < 256)
        histo_private[threadIdx.x] = 0;

    __syncthreads();
}
```



# Build Private Histogram

- Use private copies of the histo[] array to compute

```
int i = threadIdx.x + blockIdx.x * blockDim.x;

// stride is total number of threads
int stride = blockDim.x * gridDim.x;

while (i < size) {
    atomicAdd( &(private_histo[buffer[i]]), 1);
    i += stride;
}
```

# Build Final Histogram

- Copy from the histo[] arrays from each thread block to global memory

```
// wait for all other threads in the block to finish  
__syncthreads();
```

```
if (threadIdx.x < 256)  
    atomicAdd( &(histo[threadIdx.x]),  
              private_histo[threadIdx.x] );  
}
```

# More on Privatization

- Privatization is a powerful and frequently used techniques for parallelizing applications
- The operation needs to be associative and commutative
  - Histogram add operation is associative and commutative
- The histogram size needs to be small
  - Fits into shared memory
- What if the histogram is too large to privatize?

Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

**ANY MORE QUESTIONS**  
**READ CHAPTER 9**

# Problem Solving

- Q: Suppose a processor supports atomic operations in L2 cache. Assume that each atomic operation takes 5ns to complete in L2 cache and 120ns to complete in DRAM. The kernel performs 20 floating-point operations per atomic operation, and a floating-point operation takes 1ns. Assume the time of L2 atomic operations, DRAM atomic operations, and floating-point operations in each thread do not overlap. The floating-point throughput of the kernel execution is 0.2424 GFLOPS, and every thread in a block performs 5 atomic operations and 100 floating-point operations. What percent of the atomic operations happened in L2 cache?
- A:
  - 50%

# Problem Solving - Explanation

- The total time taken by a thread for all operations, assuming all atomic operations are in L2 cache, is:  $5 \text{ atomic operations} * 5\text{ns/atomic operation} + 100 \text{ floating-point operations} * 1\text{ns/floating-point operation} = 25\text{ns} + 100\text{ns} = 125\text{ns}$
- Similarly, the total time taken by a thread for all operations, assuming all atomic operations are in DRAM, is:  $5 \text{ atomic operations} * 120\text{ns/atomic operation} + 100 \text{ floating-point operations} * 1\text{ns/floating-point operation} = 600\text{ns} + 100\text{ns} = 700\text{ns}$
- Let's denote the percentage of atomic operations that happened in DRAM as  $x$ . Therefore, the percentage of atomic operations that happened in L2 cache is  $1 - x$  and the total execution time of a thread is:  $(1 - x) * 125\text{ns/thread} + x * 700\text{ns/thread}$
- Given that the floating-point throughput of the kernel execution is 0.2424 GFLOPS, we can calculate the total time taken for all operations in a thread (1 GFLOP = 1 billion floating-point operations):  $0.2424 \text{ GFLOPS} = 0.2424 * 1 \text{ billion floating-point operations/second} = 242.4 \text{ million floating-point operations/second}$ .
- Since every thread performs 100 floating-point operations, the number of threads that can be executed per second is:  $242.4 \text{ million threads/second} = 242.4 \text{ million floating-point operations/second} / 100 \text{ floating-point operations/thread} = 2.424 \text{ million threads/second}$ .
- Therefore, the total time taken for all operations in a thread is:  $1 / 2.424 \text{ million threads/second} = 0.0000004125 \text{ seconds/thread} = 412.5\text{ns/thread}$ .
- Thus:  $(1 - x) * 125\text{ns/thread} + x * 700\text{ns/thread} = 412.5\text{ns/thread}$ . Solving this for  $x$  gives us 50%.