

KERNEL PROFILING USING NSIGHT



GPA - 2

FRANK, HERMAN, HASAN,

ANSWER:

How do you know your code is running fast?

How do you know your code is memory efficient?

How do you know you're fully utilizing your SMs?

ANSWER:



HOW DO WE USE NSIGHT

OVERVIEW

- Installation & Running
- Analysis
- Matrix Multiplication Example
- Performance Metrics Deep Dive

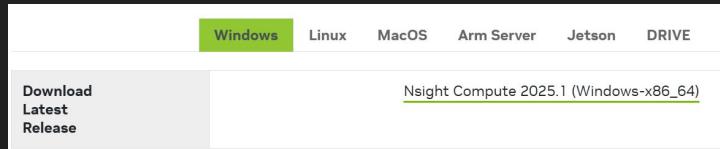
INSTALLATIO

N

INSTALLATION

1) Install Nsight Compute

- a) Go to: <https://developer.nvidia.com/tools-overview/nsight-compute/get-started>



RUNNING

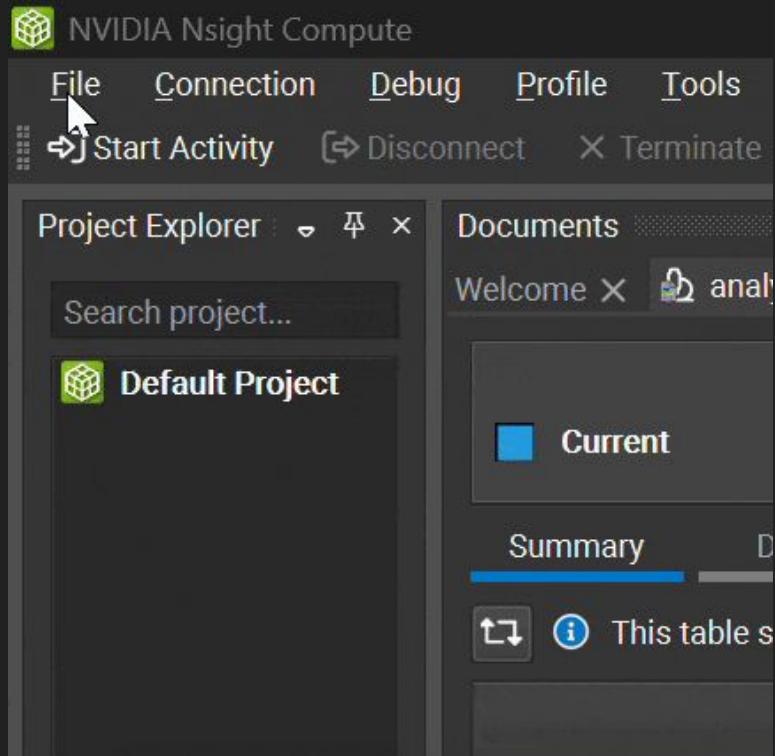
- Modify your slurm file
 - Change --constraint from --constraint="projects" to --constraint="projects,perf,nvperf"
 - Add to file: ncu --set full -f -o **outputFileName** ./**test_file** > **test_file.out**

ANALYSIS

ANALYSIS

To analyze:

- Download the .ncu-report file to your machine
- Load the file into Nsight Compute!



Here's how it should look like!

NVIDIA Nsight Compute

File Connection Debug Profile Tools Window Help

Start Activity Disconnect Terminate Profile Kernel Baselines Metric Details Launch Details

Project Explorer Documents

Welcome analysis_file_baseline.ncu.rep

Result Size Time Cycles GPU SM Frequency Process Attributes

Current 2461 - encoder_forward_kernel (192, 1, 1)x(1024, 1, 1) 7.14 us 8,946 0 - NVIDIA A40 1.25 Ghz [3392998] test_gpt2

Summary Details Source Context Comments Raw Session Compare Tools View Export

This table shows all results in the report. Use the column headers to sort the results in this report. Double-click a result to see detailed metrics. Double-click on demangled names to rename it.

ID	Estimated Speedup [%]	Function Name	Demangled Name	Duration [ms] (392.45 ms)	Runtime Improvement [ms] (0.00 ms)	Compute Throughput [%]	Memory Throughput [%]	# Registers [register/thread]
0	0.00	encoder_forward...	encoder_forward...	0.01	0.00	17.27	18.67	16
1	0.00	layernorm_forward...	layernorm_forward...	0.81	0.00	0.88	1.03	46
2	0.00	matmul_forward_k...	matmul_forward_k...	5.05	0.00	10.27	84.74	38
3	0.00	permute_kernel...	permute_kernel.co...	0.01	0.00	18.61	32.86	22
4	0.00	compute_attention...	compute_attention...	0.16	0.00	8.88	72.28	36
5	0.00	softmax_forward_k...	softmax_forward_k...	0.17	0.00	0.31	4.81	40
6	0.00	attention_weighted...	attention_weighted...	0.03	0.00	45.73	45.73	29
7	0.00	unpermute_kernel...	unpermute_kernel...	0.01	0.00	24.24	14.34	21
8	0.00	matmul_forward_k...	matmul_forward_k...	1.99	0.00	8.69	71.54	38
9	0.00	residual_forward_k...	residual_forward_k...	0.01	0.00	8.09	44.38	16
10	0.00	layernorm_forward...	layernorm_forward...	0.83	0.00	0.08	1.08	40
11	0.00	matmul_forward_k...	matmul_forward_k...	7.00	0.00	9.86	81.38	38
12	0.00	gelu_forward_kernel	gelu_forward.kerne...	0.01	0.00	22.88	44.41	16
13	0.00	matmul_forward_k...	matmul_forward_k...	7.93	0.00	8.69	71.69	38
14	0.00	residual_forward_k...	residual_forward_k...	0.01	0.00	7.86	42.75	16
15	0.00	layernorm_forward...	layernorm_forward...	0.81	0.00	0.08	1.05	40
16	0.00	matmul_forward_k...	matmul_forward_k...	4.95	0.00	10.45	86.21	38
17	0.00	permute_kernel...	permute_kernel.co...	0.01	0.00	18.78	32.23	22
18	0.00	compute_attention...	compute_attention...	0.16	0.00	8.87	72.18	36
19	0.00	softmax_forward_k...	softmax_forward_k...	0.17	0.00	0.31	4.83	40
20	0.00	attention_weighted...	attention_weighted...	0.03	0.00	45.22	45.22	29

No quantifiable performance optimization opportunities were discovered for this result. See the Details page for non-quantifiable optimization advice.

Metric Details Search metrics in current report or for a cl

Open an Nsight Compute profiling report, then click a metric on the Details or Raw page. Use the search bar to lookup metrics in the focused report's current result.

Starting info:

Result	Size	Time	Cycles	GPU	SM Frequency	
<input checked="" type="checkbox"/> Current 2560 - attention_weighted_sum_kernel	(192, 1, 1)x(1024, 1, 1)	32.86 us	42,173	0 - NVIDIA A40	1.28 Ghz	[...]
Summary	Details	Source	Context	Comments	Raw	Session

SUMMAR

Shows all kernels and executions, alongside basic runtime info.

DETAIL

Detailed information on all metrics collected by Nsight Compute.

SOURC

Source code of selected kernel in Summary page.

CONTEX

(optional) shows CPU call stack at the time of kernel launch.

COMMENT

Shows comments written in profiling file.

RAW

Shows ALL metrics collected for the selected kernel and its data.

SUMMARY

ID	Estimated Speedup [%]	Function Name	Demangled Name	Duration [ms] (392.45 ms)	Runtime Improvement [ms] (0.00 ms)	Compute Throughput [%]	Memory Throughput [%]	# Registers [register/thread]
0	0.00	encoder_forward_...	encoder_forward_...	0.01	0.00	17.27	18.67	16
1	0.00	layernorm_forward...	layernorm_forward...	0.81	0.00	0.08	1.03	40
2	0.00	matmul_forward_k...	matmul_forward_k...	5.05	0.00	10.27	84.74	38
3	0.00	permute_kernel	permute_kernel.co...	0.01	0.00	18.61	32.06	22
4	0.00	compute_attention...	compute_attention...	0.16	0.00	8.88	72.20	36
5	0.00	softmax_forward_k...	softmax_forward_k...	0.17	0.00	0.31	4.81	40
6	0.00	attention_weighted...	attention_weighted...	0.03	0.00	45.73	45.73	29
7	0.00	unpermute_kernel	unpermute_kernel..	0.01	0.00	24.24	14.34	21
8	0.00	matmul_forward_k...	matmul_forward_k...	1.99	0.00	8.69	71.54	38
9	0.00	residual_forward_k...	residual_forward_k...	0.01	0.00	8.09	44.30	16
10	0.00	layernorm_forward...	layernorm_forward...	0.83	0.00	0.08	1.00	40
11	0.00	matmul_forward_k...	matmul_forward_k...	7.00	0.00	9.86	81.38	38
12	0.00	gelu_forward_kernel	gelu_forward_kerne...	0.01	0.00	22.00	44.41	16
13	0.00	matmul_forward_k...	matmul_forward_k...	7.93	0.00	8.69	71.69	38
14	0.00	residual_forward_k...	residual_forward_k...	0.01	0.00	7.86	42.75	16
15	0.00	layernorm_forward...	layernorm_forward...	0.81	0.00	0.08	1.05	40
16	0.00	matmul_forward_k...	matmul_forward_k...	4.95	0.00	10.45	86.21	38
17	0.00	permute_kernel	permute_kernel.co...	0.01	0.00	18.78	32.23	22
18	0.00	compute_attention...	compute_attention...	0.16	0.00	8.87	72.18	36
19	0.00	softmax_forward_k...	softmax_forward_k...	0.17	0.00	0.31	4.83	40
20	0.00	attention_weighted...	attention_weighted...	0.03	0.00	45.22	45.22	29

BREAKDOWN

ID

ID	▲	Esti
0		
1		
2		
3		
4		
5		
6		
7		
8		
9		
10		
11		
12		
13		
14		
15		
16		
17		
18		
19		
20		

BREAKDOWN

ESTIMATED

BREAKDOWN

FUNCTION

DEMANGLED

Function Name	Demangled Name
encoder_forward_kernel	encoder_forward_kernel(float *, const int *, const float *, const float *, int, int, int)
layernorm_forward_kernel	layernorm_forward_kernel(float *, float *, float *, const float *, const float *, const float *, int, int, int)
matmul_forward_kernel	matmul_forward_kernel(float *, const float *, const float *, const float *, int, int)
permute_kernel	permute_kernel(float *, float *, float *, const float *, int, int, int, int)
compute_attention_scores_kernel	compute_attention_scores_kernel(float *, const float *, const float *, int, int, int, int)
softmax_forward_kernel	softmax_forward_kernel(float *, float, const float *, int, int)
attention_weighted_sum_kernel	attention_weighted_sum_kernel(float *, const float *, const float *, int, int, int, int)
unpermute_kernel	unpermute_kernel(float *, float *, int, int, int, int)
matmul_forward_kernel	matmul_forward_kernel(float *, const float *, const float *, const float *, int, int)
residual_forward_kernel	residual_forward_kernel(float *, float *, float *, int)
layernorm_forward_kernel	layernorm_forward_kernel(float *, float *, float *, const float *, const float *, const float *, int, int, int)
matmul_forward_kernel	matmul_forward_kernel(float *, const float *, const float *, const float *, int, int)

BREAKDOWN

Duration [ms] (392.45 ms)	Runtime Improvement [ms] (0.00 ms)
0.01	0.00
0.81	0.00
5.05	0.00
0.01	0.00
0.16	0.00
0.17	0.00
0.03	0.00
0.01	0.00
1.99	0.00
0.01	0.00
0.83	0.00
7.00	0.00
0.01	0.00
7.93	0.00
0.01	0.00
0.81	0.00
4.95	0.00
0.01	0.00
0.16	0.00

DURATION [MS] (TOTAL

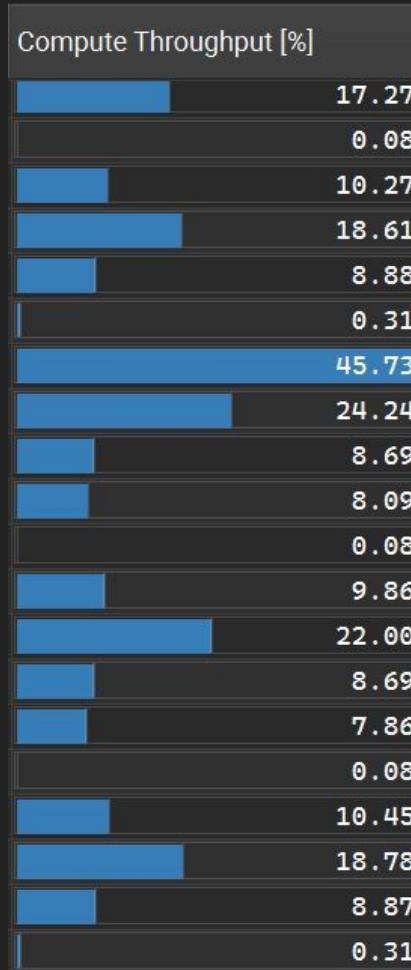
RUNTIME IMPROVEMENT [MS] (TOTAL

TIME

BREAKDOWN

COMPUTE THROUGHPUT

[%]



BREAKDOWN

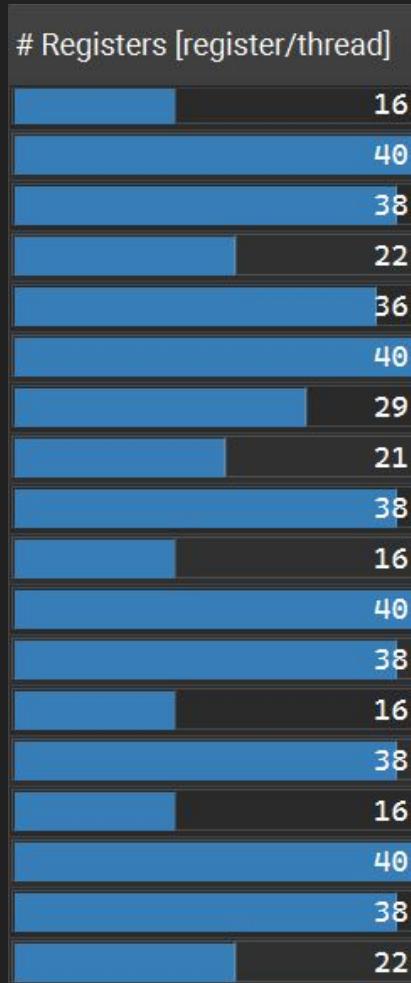
MEMORY THROUGHPUT [%]

Memory Throughput [%]
18.67
1.03
84.74
32.06
72.20
4.81
45.73
14.34
71.54
44.30
1.00
81.38
44.41
71.69
42.75
1.05
86.21
32.23
72.18
4.83

BREAKDOWN

REGISTERS

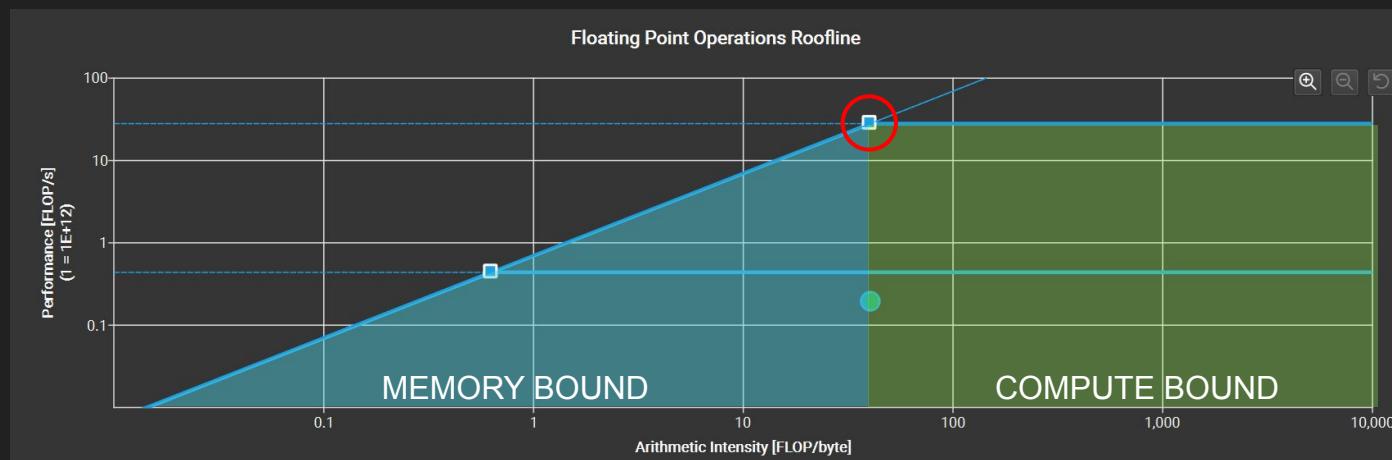
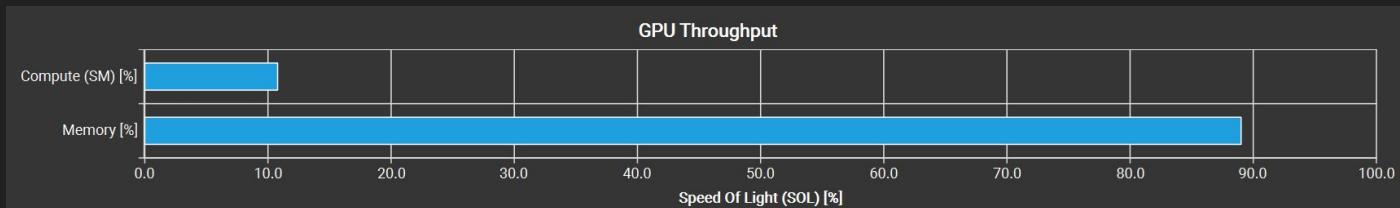
(REGISTER/THREAD)



DETAILS OVERVIEW

Speed of Light

- High-level performance at a glance
- Achieved throughput relative to theoretical maximum

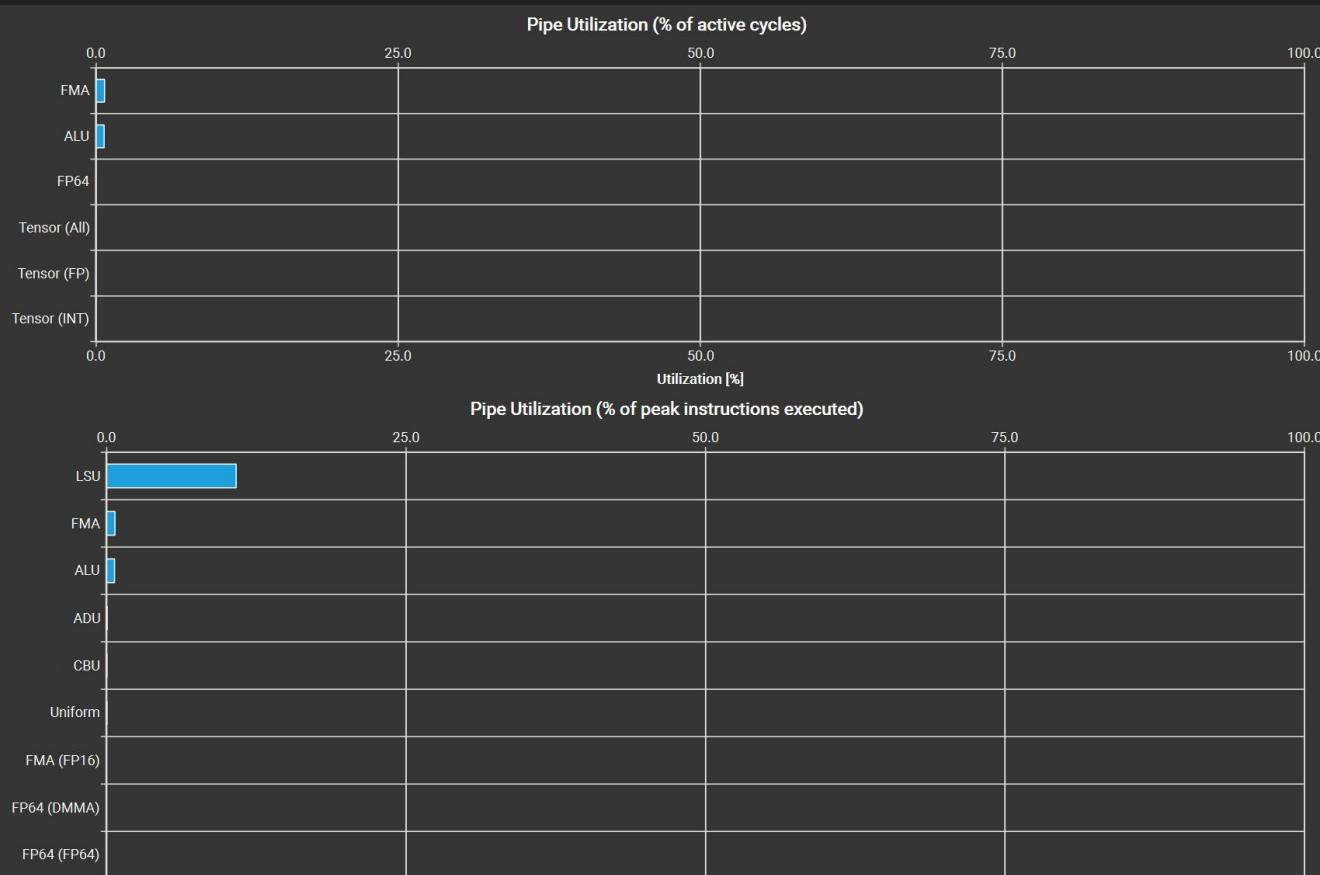


Arithmetic Intensity
= FLOP/byte

Closer to Roofline is
better!

Compute Workload Analysis

Pipe: Hardware that carries out specific instructions



1. % of Cycles when the **SM** was **active**

Low = pipe is idle, even when SM is *active*

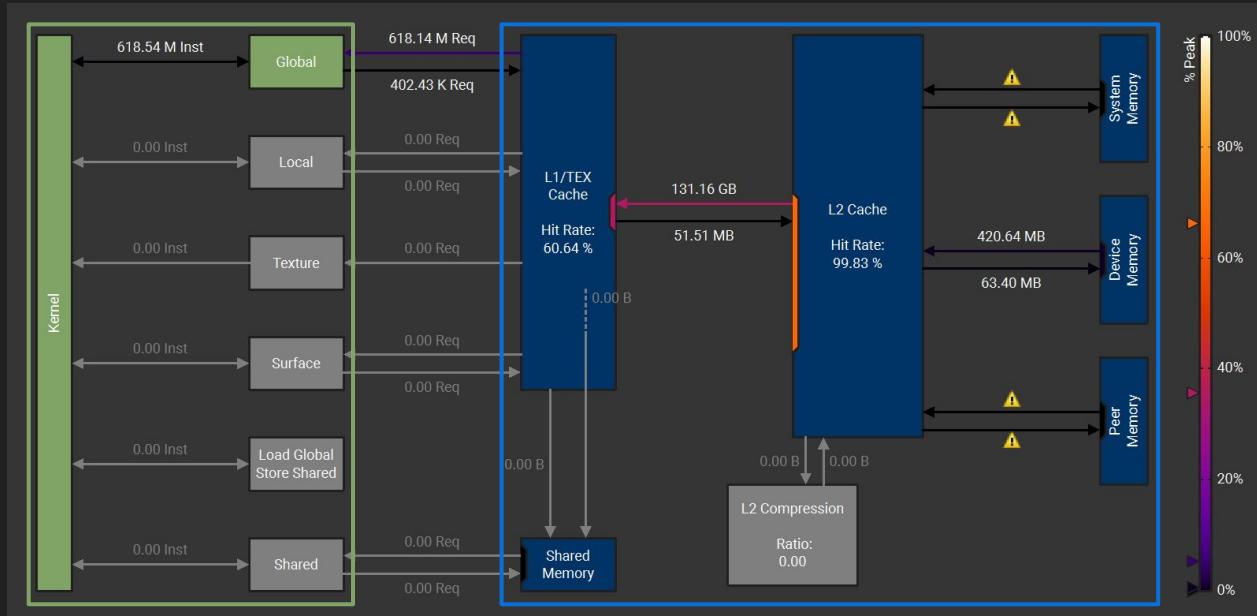
High = SMs are working when they're *active*

2. % Of theoretical peak ("Speed of light" of pipes)

Low = Memory stalls, branch divergence, etc.

High = Utilizing hardware as much as possible

Memory Workload Analysis



Logical Units

Physical Units

Almost all instructions go through L1 → L2 → DRAM

Port: Incoming and Outgoing traffic

Link: Instructions, Requests, or Data between Units

%Peak = Percentage of peak utilization of each Link/Port

Warp State Statistics

- Warp cycle
- eligible/non-eligible warp

▼ Warp State Statistics

Analysis of the states in which all warps spent cycles during the kernel execution. The warp states describe a warp's readiness or inability to issue its next instruction. The warp cycles per instruction define the latency between two consecutive instructions. The higher the value, the more warp parallelism is required to hide this latency. For each warp state, the chart shows the average number of cycles spent in that state per issued instruction. Stalls are not always impacting the overall performance nor are they completely avoidable. Only focus on stall reasons if the schedulers fail to issue every cycle. When executing a kernel with mixed library and user code, these metrics show the combined values.

Warp Cycles Per Issued Instruction [cycle]	326.09	Avg. Active Threads Per Warp	32
Warp Cycles Per Executed Instruction [cycle]	326.10	Avg. Not Predicated Off Threads Per Warp	31.90

lg_throttle On average, each warp of this kernel spends 308.5 cycles being stalled waiting for the local/global instruction queue to be not full. This represents about 94.6% of the total average of 326.1 cycles between issuing two instructions. Typically this stall occurs only when executing local or global memory instructions extremely frequently. If applicable, consider combining multiple lower-width memory operations into fewer wider memory operations and try interleaving memory operations and math instructions.

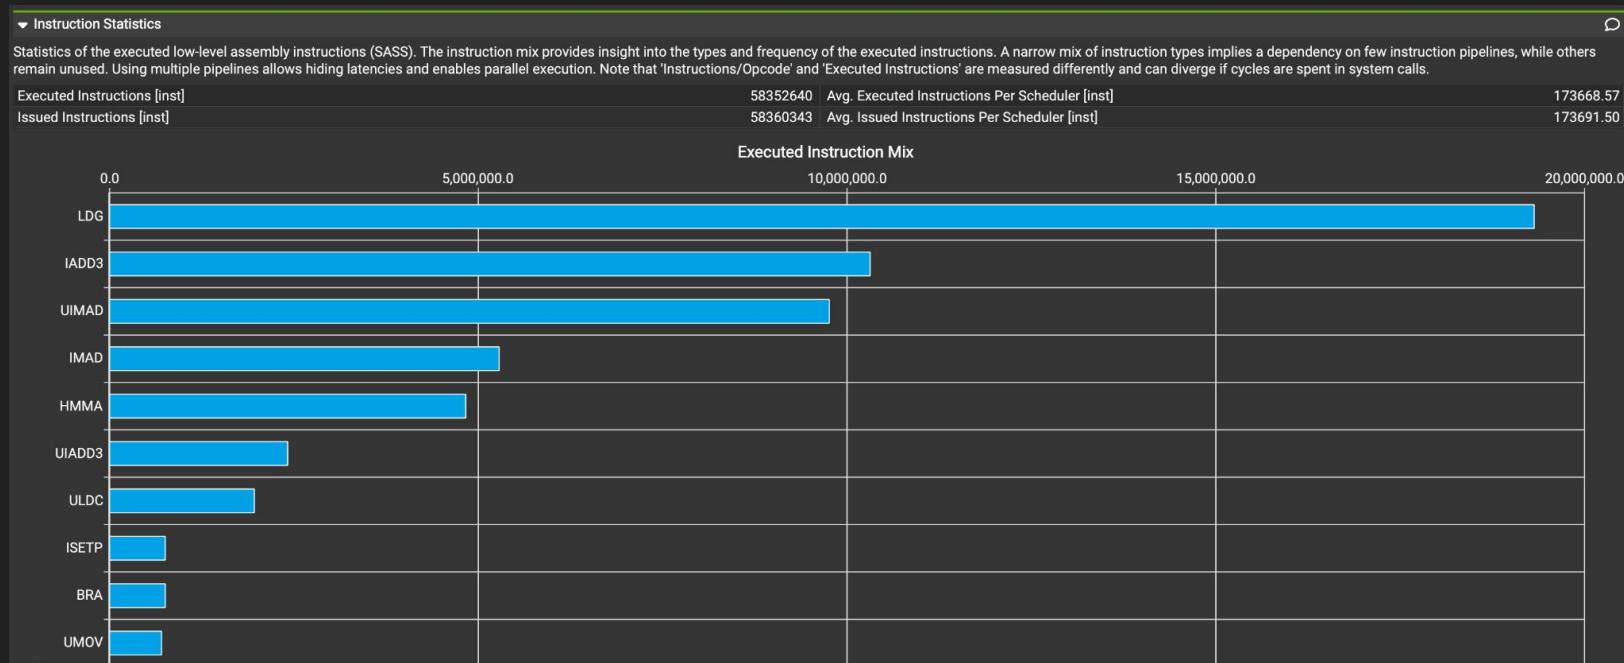
ⓘ Warp Stall Check the [Source Counters](#) section for the top stall locations in your source based on sampling data. The [Kernel Profiling Guide](#) provides more details on each stall reason.

Warp State (All Cycles)

The chart displays the average number of cycles spent in different warp states per issued instruction. The y-axis lists the stall reasons, and the x-axis shows the cycle count from 0.0 to 400.0. The bars are blue.

Warp State	Avg. Cycles Per Issued Instruction
Stall LG Throttle	308.5
Stall Long Scoreboard	~10
Stall Not Selected	~5
Stall Wait	~2
Selected	~1

Instructions Statistics



Launch Statistics & Occupancy

Actual warps active per SM / Max warps active per SM

Launch Statistics		
Summary of the configuration used to launch the kernel. The launch configuration defines the size of the kernel grid, the division of the grid into blocks, and the GPU resources needed to execute the kernel. Choosing an efficient launch configuration maximizes device utilization.		
Grid Size	12800	Function Cache Configuration
Registers Per Thread [register/thread]	38	Static Shared Memory Per Block [byte/block]
Block Size	1024	Dynamic Shared Memory Per Block [byte/block]
Threads [thread]	13107200	Driver Shared Memory Per Block [Kbyte/block]
Waves Per SM	152.38	Shared Memory Configuration Size [Kbyte]
Occupancy		
Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.		
Theoretical Occupancy [%]	66.67	Block Limit Registers [block]
Theoretical Active Warps per SM [warp]	32	Block Limit Shared Mem [block]
Achieved Occupancy [%]	65.70	Block Limit Warps [block]
Achieved Active Warps Per SM [warp]	31.54	Block Limit SM [block]

Block Limit Registers: total registers per SM / (registers per thread×threads per block)

Block Limit Shared Mem: Shared mem per SM / Shared mem per block

Block Limit Warps: Max warps per SM / Warps per block

Block Limit SM: hard architecture limit 16 blocks per SM

Source Counters

Summary Details Source Context Comments Raw Session Compare Tools View Export All

Source metrics, including branch efficiency and sampled warp stall reasons. Warp Stall Sampling metrics are periodically sampled over the kernel runtime. They indicate when warps were stalled and couldn't be scheduled. See the documentation for a description of all stall reasons. Only focus on stalls if the schedulers fail to issue every cycle.

Branch Instructions [inst]	22945792	Branch Efficiency [%]	100
Branch Instructions Ratio [%]	0.02	Avg. Divergent Branches	0

⚠️ Uncoalesced Global Accesses This kernel has uncoalesced global accesses resulting in a total of 8653897728 excessive sectors (85% of the total 10200846336 sectors). Check the L2 Theoretical Sectors Global Excessive table for the primary source locations. The [CUDA Programming Guide](#) had additional information on reducing uncoalesced device memory accesses.

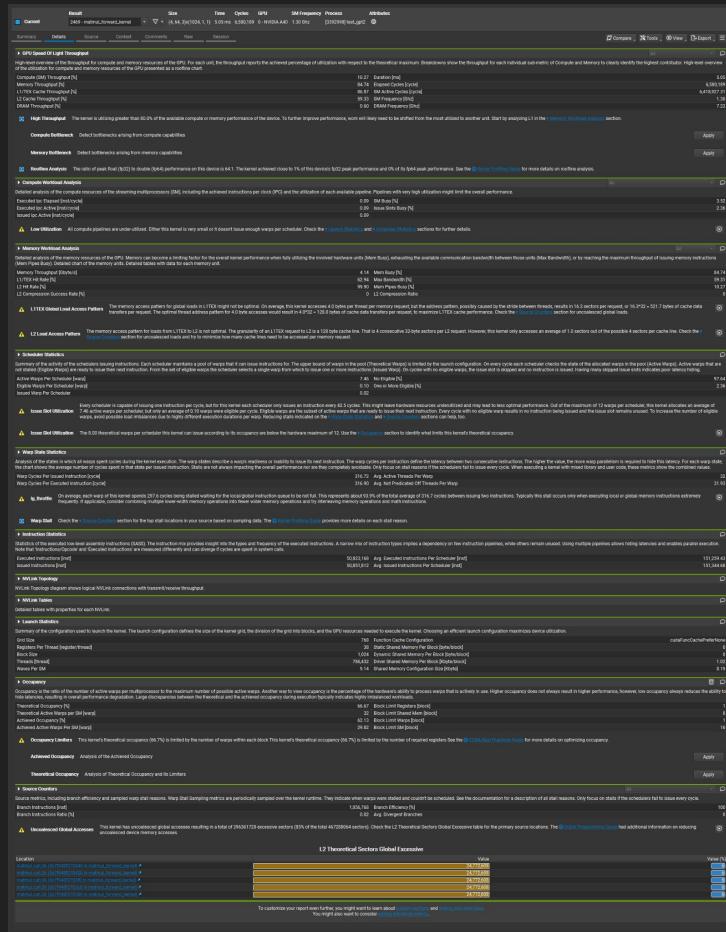
Location	Value	Value (%)
matmul.cuh:26 (0x7f948f270440 in matmul_forward_kernel) ↗	540,868,608	6
matmul.cuh:26 (0x7f948f270420 in matmul_forward_kernel) ↗	540,868,608	6
matmul.cuh:26 (0x7f948f2703f0 in matmul_forward_kernel) ↗	540,868,608	6
matmul.cuh:26 (0x7f948f2703c0 in matmul_forward_kernel) ↗	540,868,608	6
matmul.cuh:26 (0x7f948f270380 in matmul_forward_kernel) ↗	540,868,608	6

Location	Value	Value (%)
matmul.cuh:26 (0x7f948f270210 in matmul_forward_kernel) ↗	518,478	14
matmul.cuh:26 (0x7f948f2703c0 in matmul_forward_kernel) ↗	206,982	6
matmul.cuh:26 (0x7f948f270380 in matmul_forward_kernel) ↗	152,536	4
matmul.cuh:26 (0x7f948f270350 in matmul_forward_kernel) ↗	150,450	4
matmul.cuh:26 (0x7f948f2703f0 in matmul_forward_kernel) ↗	148,065	4

EXAMPLE

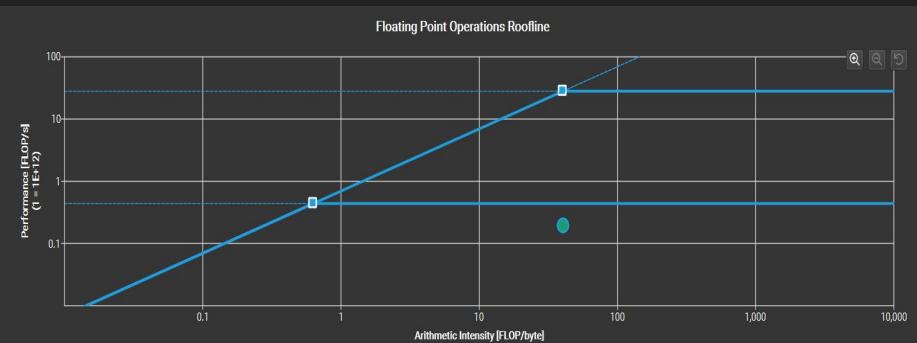
Simple example: Matmul

- $(4, 64, 3) * (1024, 1, 1)$
- Batched matrix multiplication kernel for GPT-2's forward propagation
- It performs linear transformation for each token embedding.
- The whole ncu report looks like this
- It's a lot, but no need to panic



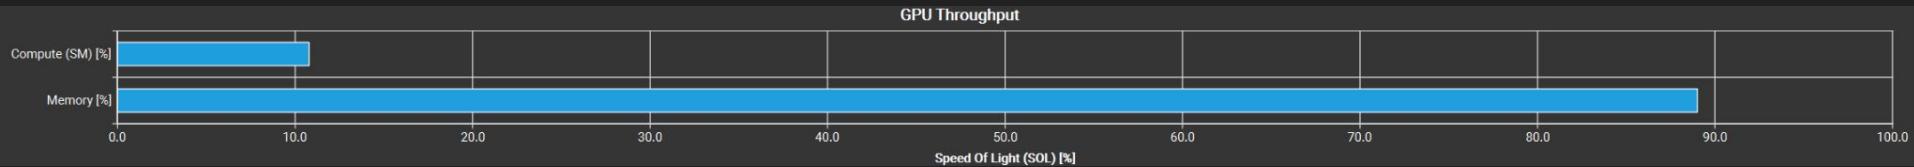
Why is that kernel “BAD”

- **Memory Bottleneck**
- Very poor memory coalescing
- Global memory inefficiency/overuse
- Abysmally low compute throughput
- Warps stall constantly



High-level overview of the throughput for compute and memory resources of the GPU. For each unit, the throughput reports the achieved percentage of utilization with respect to the maximum theoretical throughput.

Compute (SM) Throughput [%]	10.79
Memory Throughput [%]	89.01
L1/TEX Cache Throughput [%]	89.19
L2 Cache Throughput [%]	66.30
DRAM Throughput [%]	0.66



Why is that kernel “BAD”

- Memory Bottleneck
- **Very poor memory coalescing**
- Global memory inefficiency/overuse
- Abysmally low compute throughput
- Warps stall constantly

 **Uncoalesced Global Accesses** This kernel has uncoalesced global accesses resulting in a total of 8653897728 excessive sectors (85% of the total 10200846336 sectors). Check the L2 Theoretical Sectors Global Excessive table for the primary source locations. The [CUDA Programming Guide](#) had additional information on reducing uncoalesced device memory accesses.

The following table lists the metrics that are key performance indicators:

Metric Name	Value	Guidance
derived_memory_l2_theoretical_sectors_global_excessive	8.6539e+09	memory_l2_theoretical_sectors_global > memory_l2_theoretical_sectors_global_ideal

`derived__memory_l2_theoretical_sectors_global_excessive`

Excessive theoretical number of sectors requested in L2 from global memory instructions, because not all not predicated-off threads performed the operation.

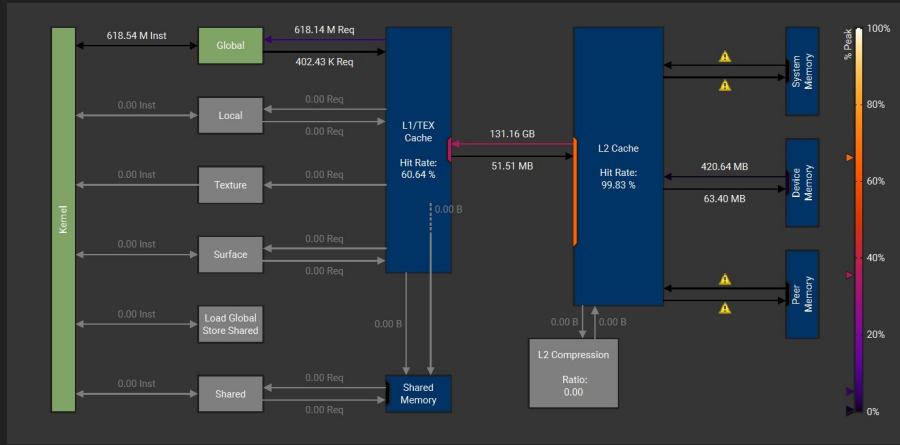
Why is that kernel “BAD”

- Memory Bottleneck
- Very poor memory coalescing
- **Global memory inefficiency/overuse**
- Abysmally low compute throughput
- Warps stall constantly

▼ Memory Workload Analysis

Detailed analysis of the memory resources of the GPU. Memory can become a limiting factor for the overall kernel performance, either by exceeding the bandwidth between those units (Max Bandwidth), or by reaching the maximum throughput of issuing memory instructions (Max Throughput).

Memory Throughput [Gbyte/s]	4.62
L1/TEX Hit Rate [%]	60.64
L2 Hit Rate [%]	99.83
L2 Compression Success Rate [%]	0

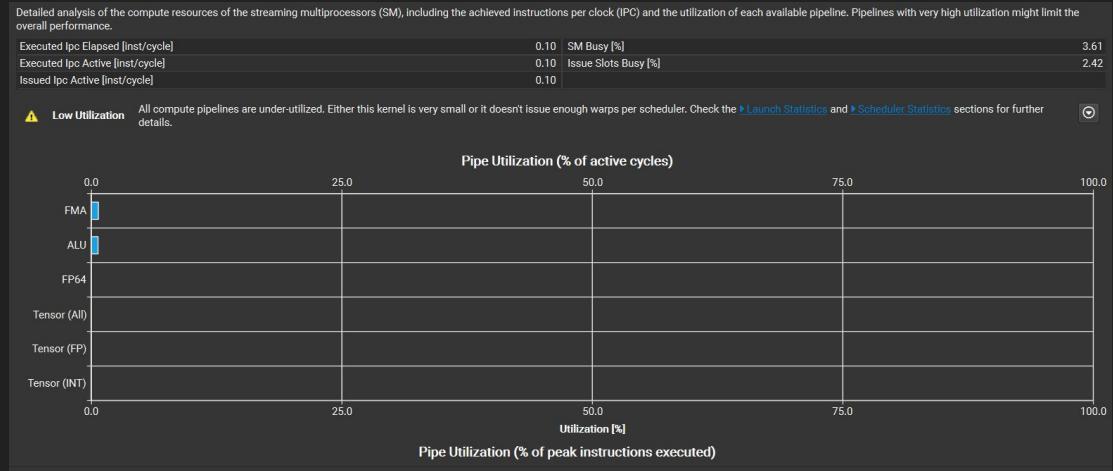


SPECIFICATIONS

GPU architecture	NVIDIA Ampere architecture
GPU memory	48 GB GDDR6 with ECC
Memory bandwidth	696 GB/s
Interconnect interface	NVIDIA® NVLink® 112.5 GB/s (bidirectional) ³ PCIe Gen4: 64GB/s
NVIDIA Ampere architecture-based CUDA Cores	10,752
NVIDIA second-generation RT Cores	84
NVIDIA third-generation Tensor Cores	336
Peak FP32 TFLOPS (non-Tensor)	37.4

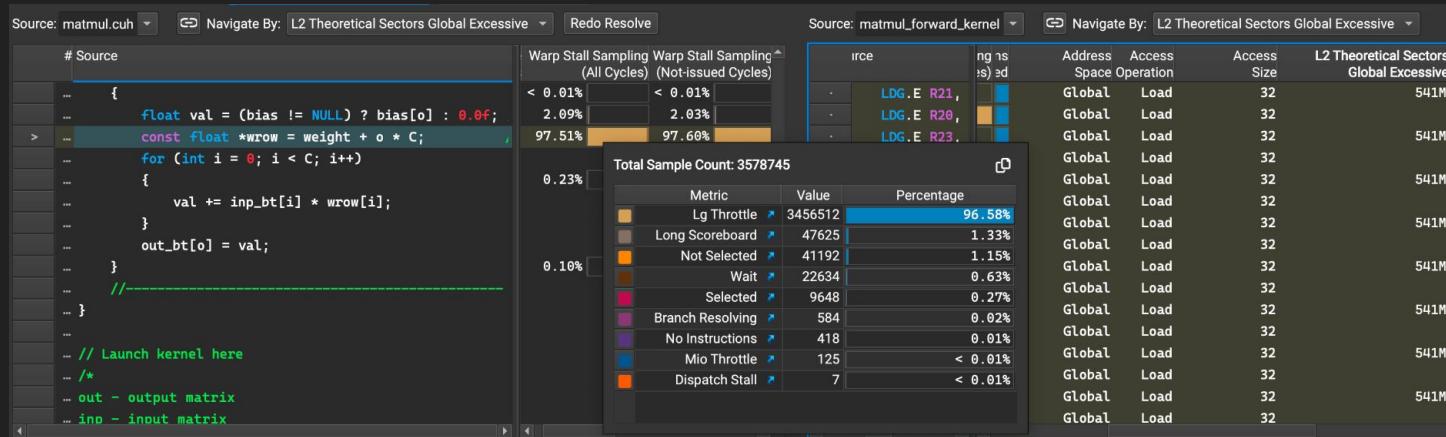
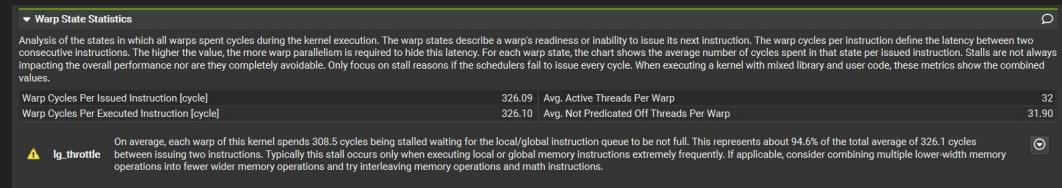
Why is that kernel “BAD”

- Memory Bottleneck
- Very poor memory coalescing
- Global memory inefficiency/overuse
- **Abysmally low compute throughput**
- Warps stall constantly



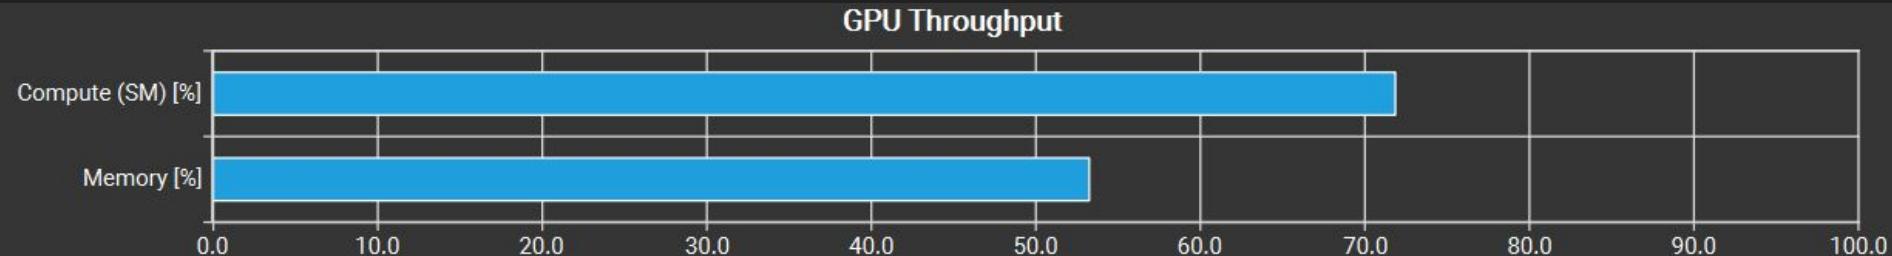
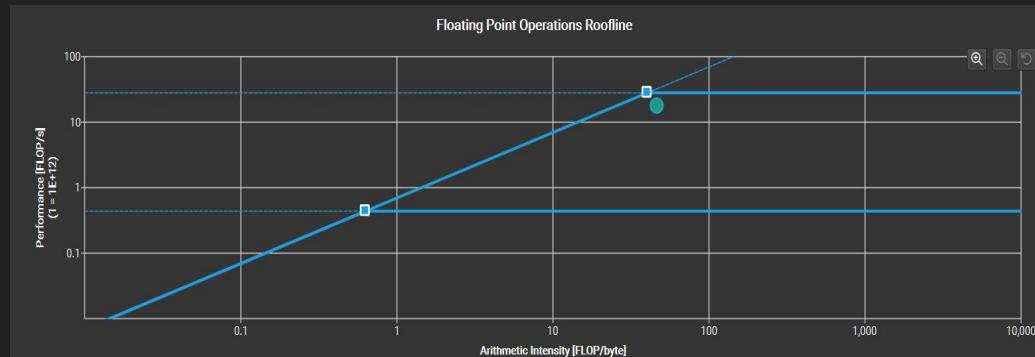
Why is that kernel “BAD”

- Memory Bottleneck
 - Very poor memory coalescing
 - Global memory inefficiency/overuse
 - Abysmally low compute throughput
 - **Warps stall constantly**



Improvement: cuBLAS optimization

- Faster: 104.77ms -> 2.29ms
- **No More Memory Bottleneck**
- Improved memory coalescing
- Shared Memory Utilization
- Satisfying compute throughput
- Warps issue more instructions



Improvement: cuBLAS optimization

- Faster: 104.77ms -> 2.29ms
- No More Memory Bottleneck
- **Improved memory coalescing**
- Shared Memory Utilization
- Satisfying compute throughput
- Warps issue more instructions

 **Uncoalesced Shared Accesses**

This kernel has uncoalesced shared accesses resulting in a total of 804864 excessive wavefronts (1% of the total 89754912 wavefronts). Check the L1 Wavefronts Shared Excessive table for the primary source locations. The [CUDA Best Practices Guide](#) has an example on optimizing shared memory accesses.

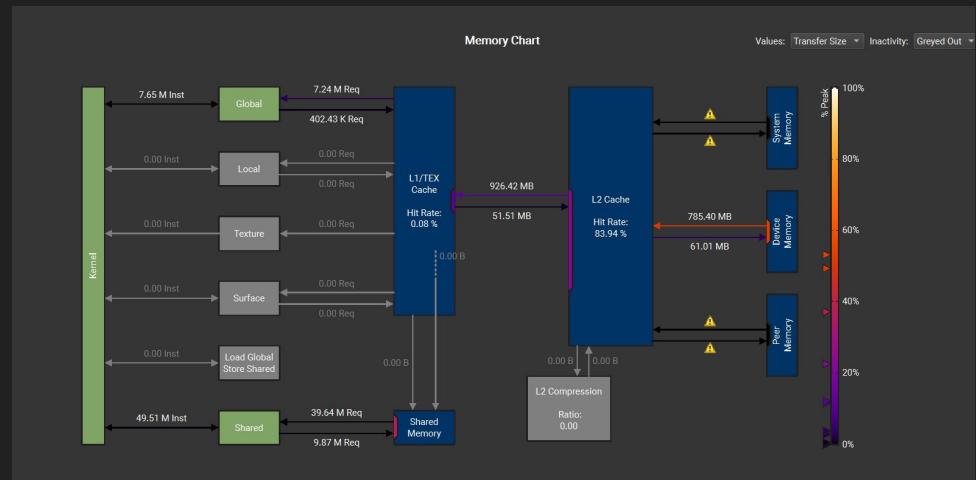
The following table lists the metrics that are key performance indicators:

Metric Name	Value	Guidance
derived_memory_l1_wavefronts_shared_excessive	804864	memory_l1_wavefronts_shared > memory_l1_wavefronts_shared_ideal

Still minor issues, but overall much better

Improvement: cuBLAS optimization

- Faster: 104.77ms -> 2.29ms
- No More Memory Bottleneck
- Improved memory coalescing
- **Shared Memory Utilization**
- Satisfying compute throughput
- Warps issue more instructions



Improvement: cuBLAS optimization

- Faster: 104.77ms -> 2.29ms
- No More Memory Bottleneck
- Improved memory coalescing
- Shared Memory Utilization
- **Satisfying compute throughput**
- Warps issue more instructions

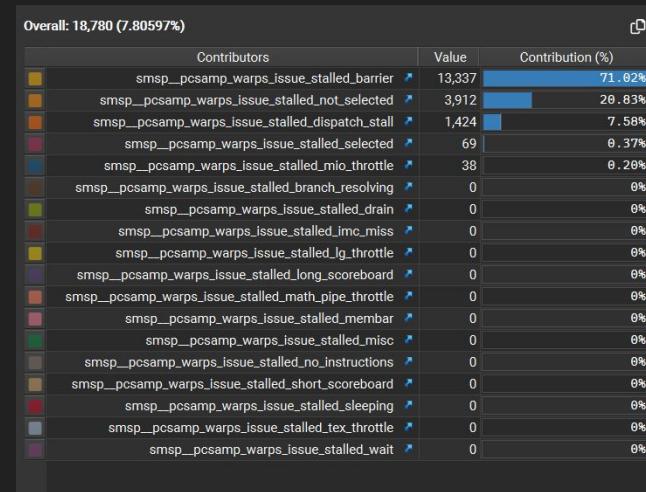


Detailed analysis of the compute resources of the streaming multiprocessors (SM), including the achieved instructions per clock (IPC) and the utilization of each available pipeline. Pipelines with very high utilization might limit the overall performance.			
Executed Ipc Elapsed [inst/cycle]	0.10	SM Busy [%]	3.61
Executed Ipc Active [inst/cycle]	0.10	Issue Slots Busy [%]	2.42
Issued Ipc Active [inst/cycle]	0.10		

Detailed analysis of the compute resources of the streaming multiprocessors (SM), including the achieved instructions per clock (IPC) and the utilization of each available pipeline. Pipelines with very high utilization might limit the overall performance.			
Executed Ipc Elapsed [inst/cycle]	2.87	SM Busy [%]	73.02
Executed Ipc Active [inst/cycle]	2.92	Issue Slots Busy [%]	73.02
Issued Ipc Active [inst/cycle]	2.92		

Improvement: cuBLAS optimization

- Faster: 104.77ms -> 2.29ms
- No More Memory Bottleneck
- Improved memory coalescing
- Shared Memory Utilization
- Satisfying compute throughput
- **Warps issue more instructions**



Warp Cycles Per Issued Instruction [cycle]	326.09	Avg. Active Threads Per Warp	32
Warp Cycles Per Executed Instruction [cycle]	326.10	Avg. Not Predicated Off Threads Per Warp	31.90

Warp Cycles Per Issued Instruction [cycle]	5.36	Avg. Active Threads Per Warp	32
Warp Cycles Per Executed Instruction [cycle]	5.36	Avg. Not Predicated Off Threads Per Warp	31.76

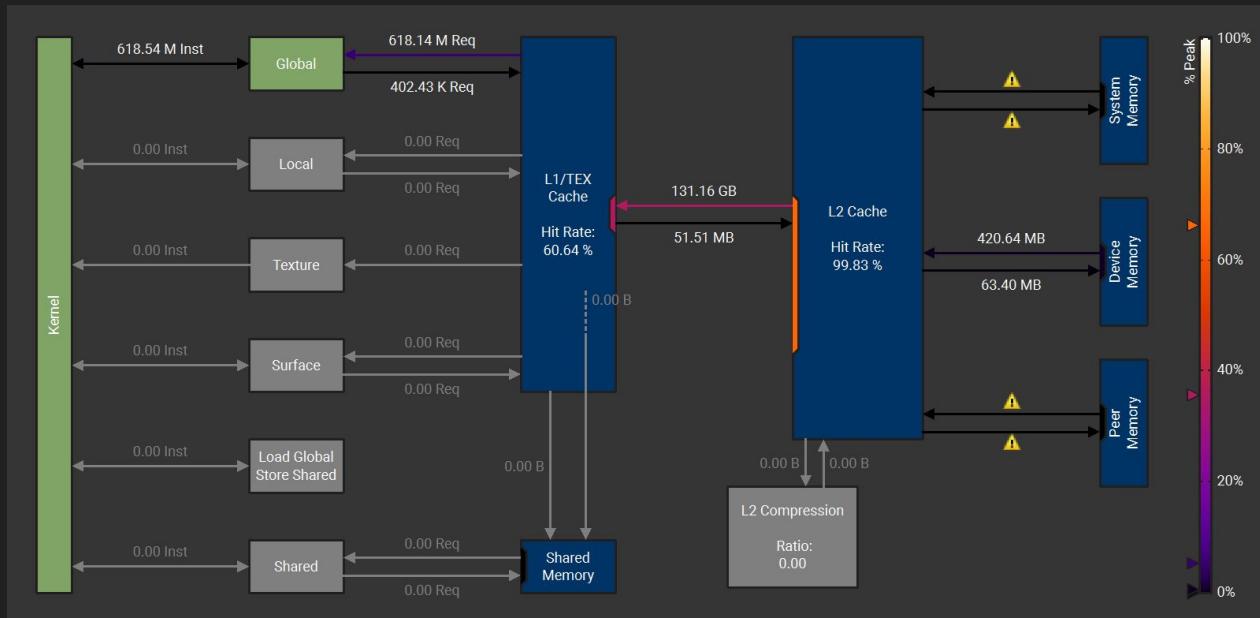
How does Nsight Compute do this?

Under non-instrumented mode of NCU:

1. Hardware performance monitoring unit
 - Tracks metrics like number and types of instructions executed, cache hits/misses, memory transactions, etc.
2. Sampling + Counters
 - Samples statistics like warps stalling, pipeline utilization
 - Counts number and types of instructions executed, number of branched instructions

PERFORMANCE METRICS DEEP DIVE

Memory Workload - Baseline Mat. Mul. (No Tiling)



618M Total Instructions

L1 Cache Under-utilized

40% (250M) Miss L1

L2 Reads cost **130 GB!**

~200 cycles per L2 load*

Cache bound

Matrices too big for cache

Repeatedly gets streamed
from DRAM → L2 → L1

For Reference:

Input Matrix: $(4 \times 64) \times 768 = 196,608$ total elements

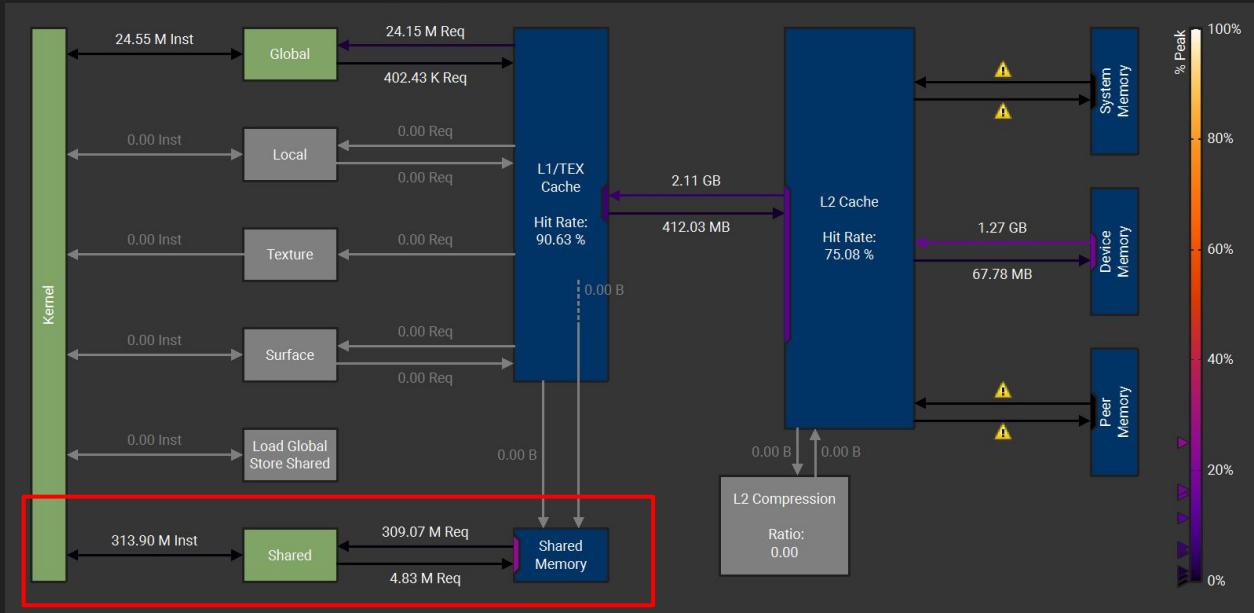
Weight Matrix: $50304 \times 768 = 38,633,472$ total elements

Total Input Size: ~155 MB

Total Available L2 Cache: 6 MB

*based on A100 GPU

Memory Workload - Register & Shared Memory Tiling



24M Global (-96%)
+ 313M Shared
= 337M (-45%) Total Instr.

L2 Cache Not Stressed
-96% less Global reads
90% L1 Hit Rate
→ L2 bandwidth freed up

T = 64: 64x Reuse with shared memory:
309.9 load / 4.8 store = 64!

For Reference:

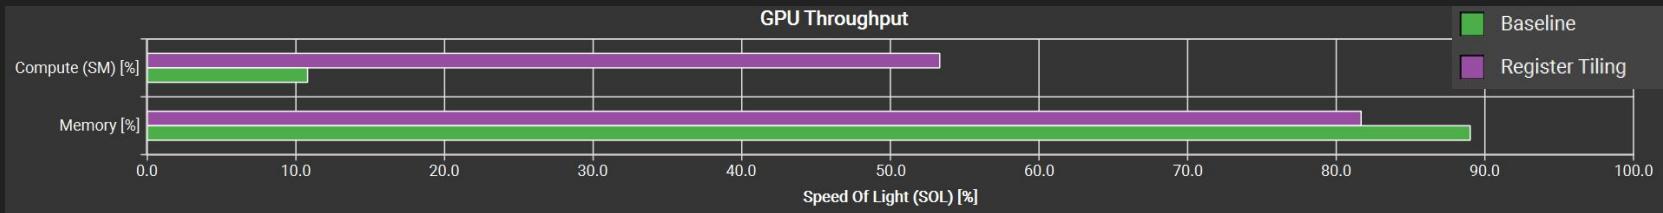
Input Matrix: $(4 \times 64) \times 768 = 196,608$ total elements

Weight Matrix: $50304 \times 768 = 38,633,472$ total elements

Total Input Size: ~155 MB

Total Available L2 Cache: 6 MB

Memory Workload - Register & Shared Memory Tiling



Results

- Compute: 10.8% → 53.3%
- Memory: 89.6% → 81%
- Runtime: 104.7ms → 11.58ms

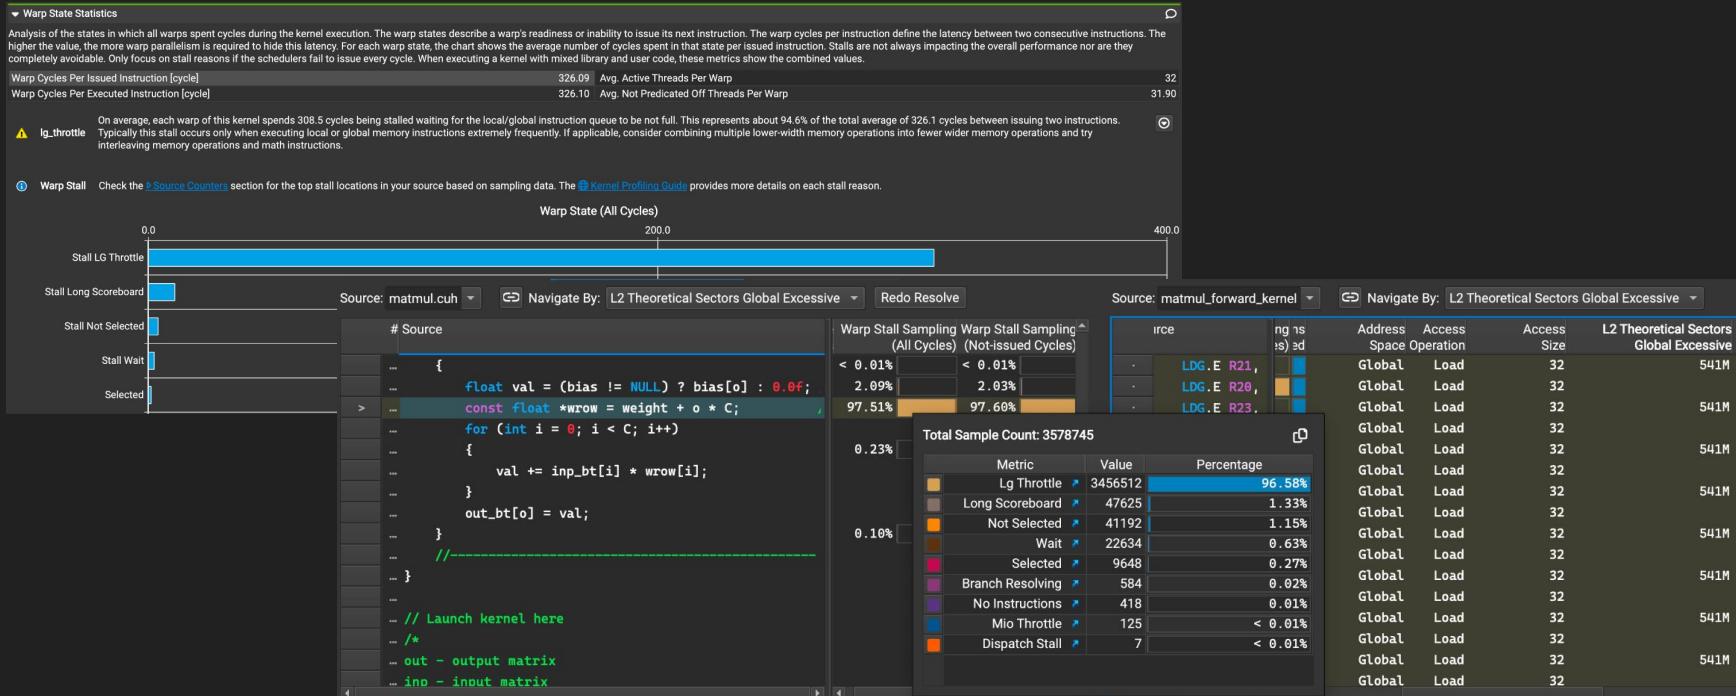
Profiling Summary

- Less cycles stalled for loading memory → increased FLOPS
 - Less Global instructions and lower L2 bandwidth
- Use of on-chip registers and shared memory → lower memory latency
 - ~5-10 cycles v.s. ~200

No Eligible 97.58%
One or More Eligible 2.42%

Baseline warp stall

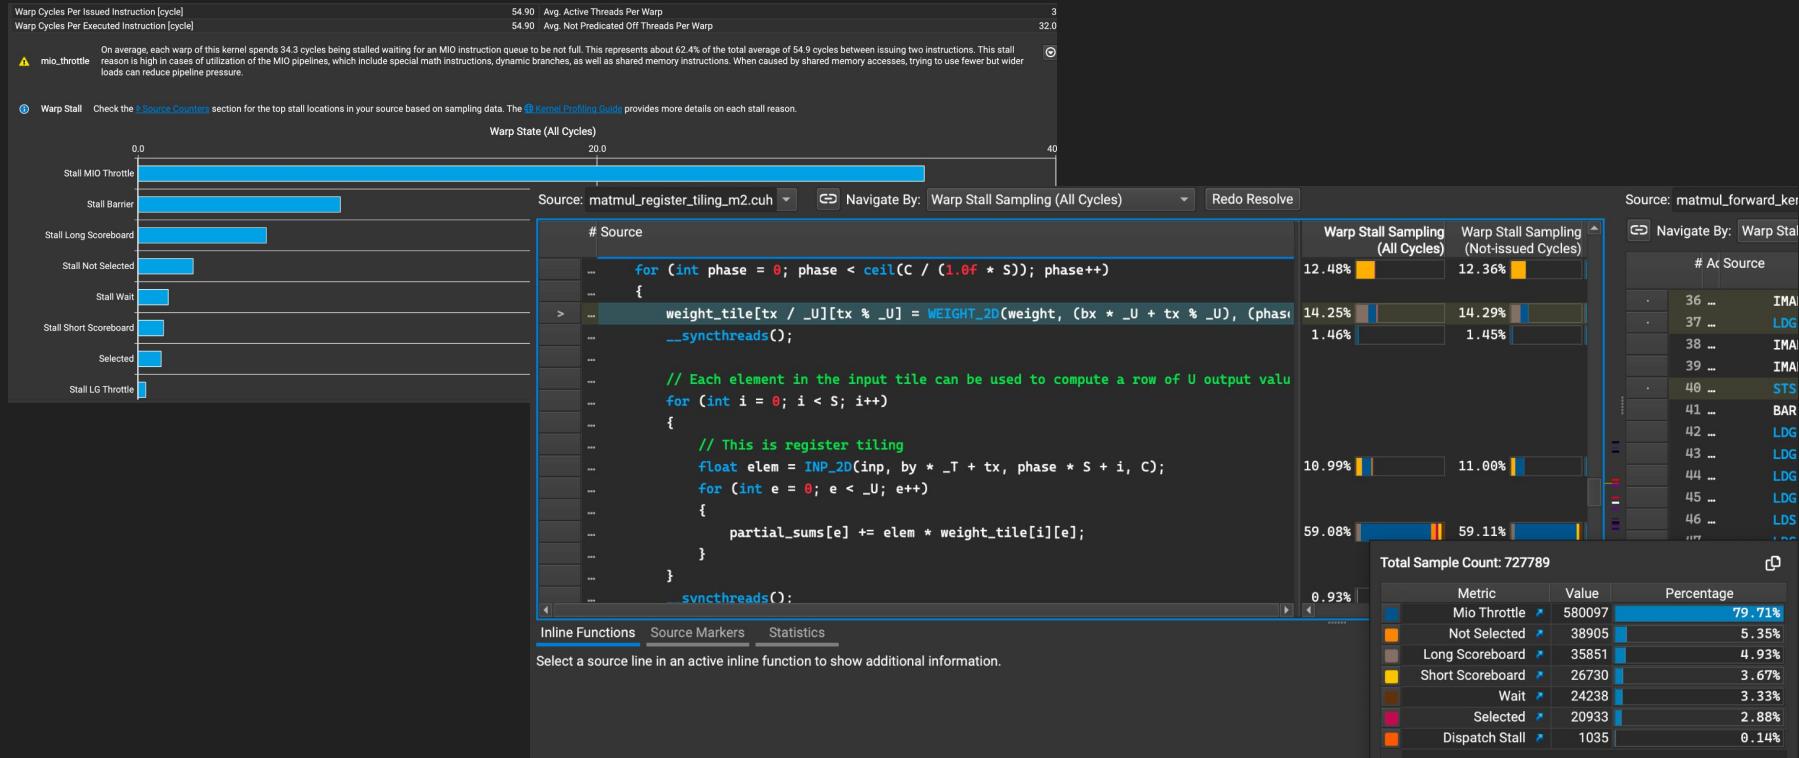
Lg Throttle – executing local or global memory instructions extremely frequently



No Eligible 85.97%
One or More Eligible 14.03%

Register tiling warp stall

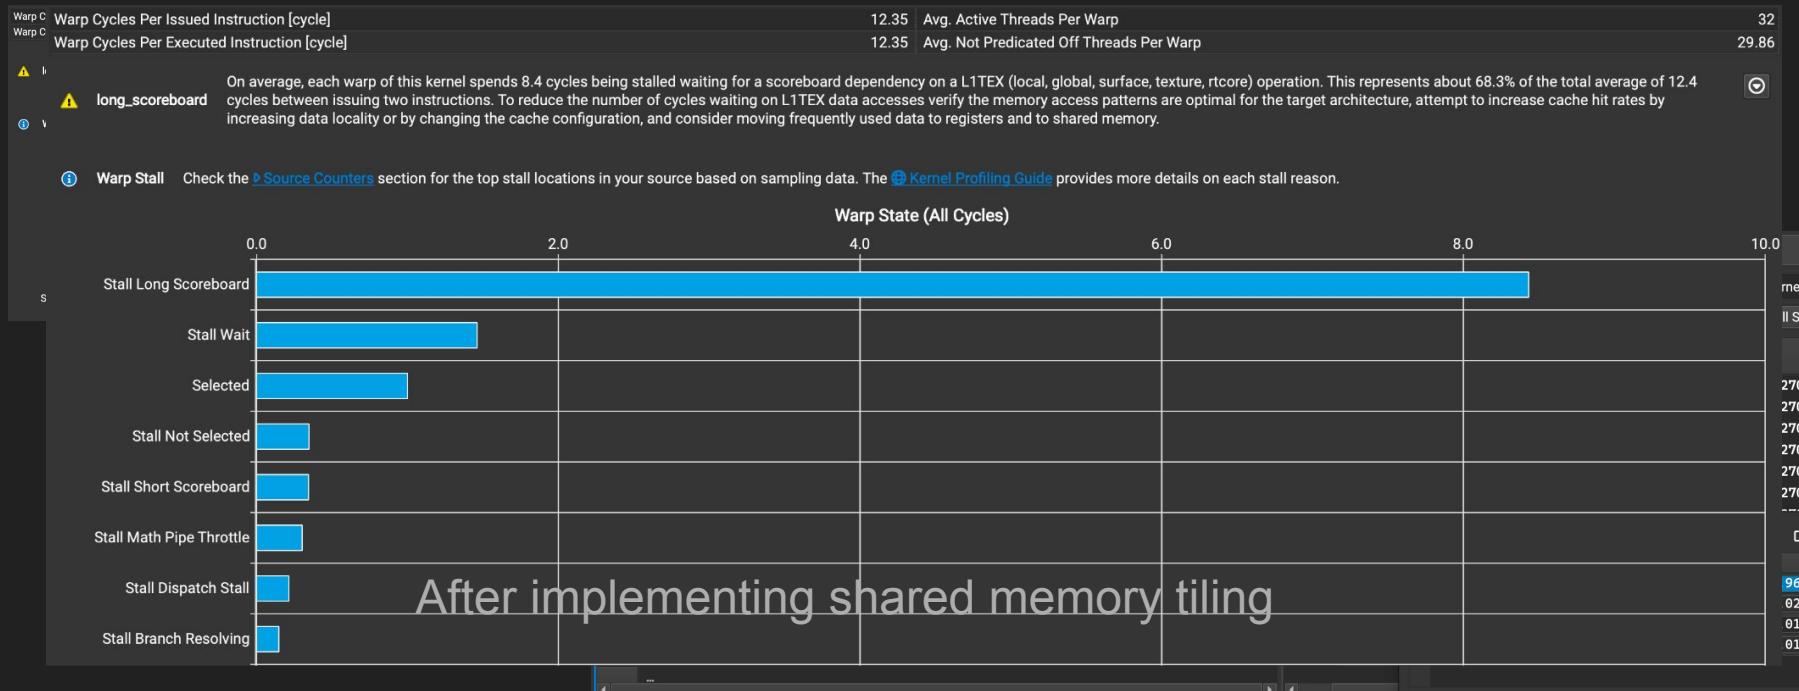
Mio_throttle – special math instructions, dynamic branches, shared memory instructions



No Eligible 65.98%
One or More Eligible 34.02%

Tensor core warp stall

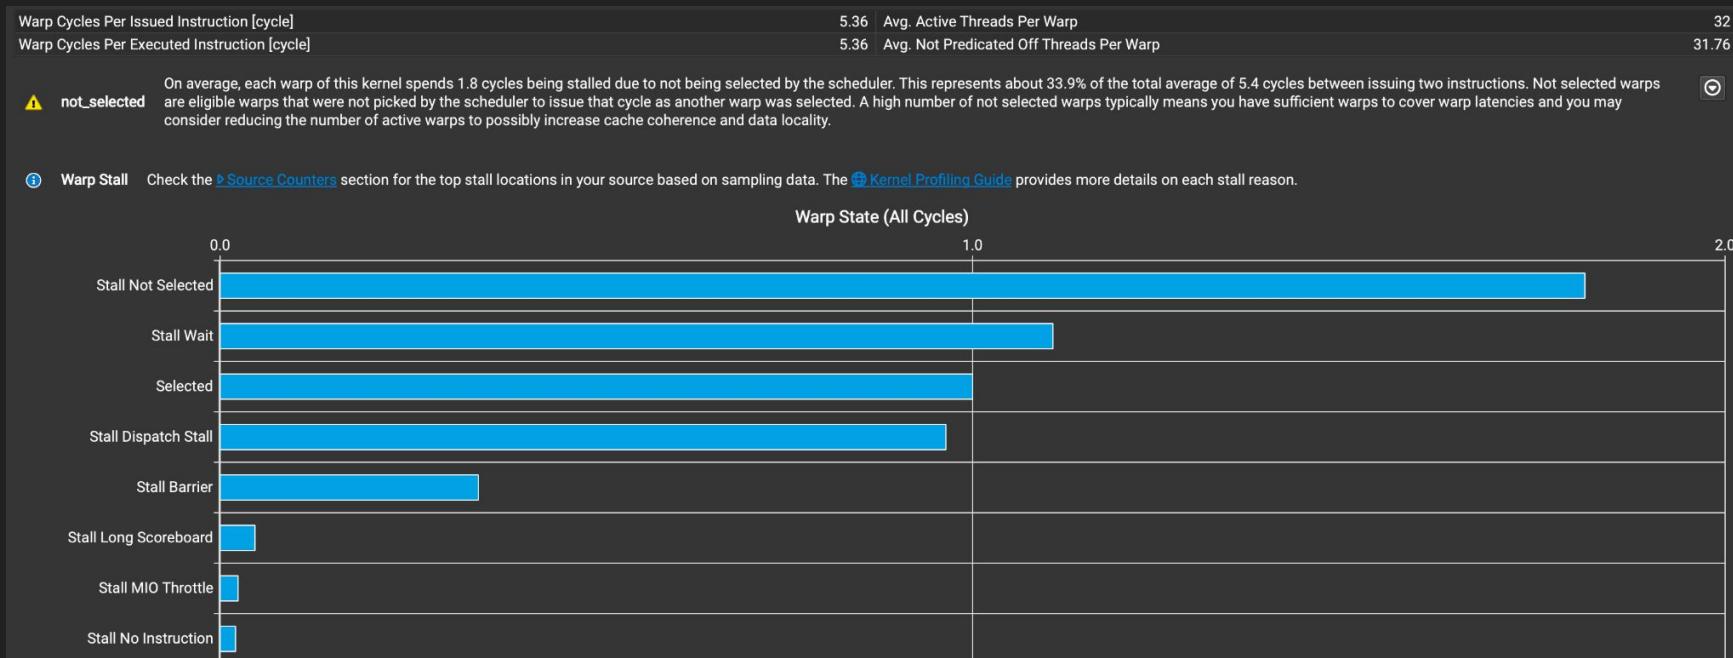
Long scoreboard – waiting on memory operations from L1TEX (global, local, or surface memory).



No Eligible 26.93%
One or More Eligible 73.07%

cuBLAS warp stall

Not selected: eligible warps that were not picked by the scheduler to issue that cycle as another warp was selected



6. Launch Statistics & Occupancy

Actual warps active per SM / Max warps active per SM

Launch Statistics		
Summary of the configuration used to launch the kernel. The launch configuration defines the size of the kernel grid, the division of the grid into blocks, and the GPU resources needed to execute the kernel. Choosing an efficient launch configuration maximizes device utilization.		
Grid Size	12800	Function Cache Configuration
Registers Per Thread [register/thread]	38	static Shared Memory Per Block [byte/block]
Block Size	1024	Dynamic Shared Memory Per Block [byte/block]
Threads [thread]	13107200	Driver Shared Memory Per Block [Kbyte/block]
Waves Per SM	152.38	Shared Memory Configuration Size [Kbyte]
Occupancy		
Occupancy is the ratio of the number of active warps per multiprocessor to the maximum number of possible active warps. Another way to view occupancy is the percentage of the hardware's ability to process warps that is actively in use. Higher occupancy does not always result in higher performance, however, low occupancy always reduces the ability to hide latencies, resulting in overall performance degradation. Large discrepancies between the theoretical and the achieved occupancy during execution typically indicates highly imbalanced workloads.		
Theoretical Occupancy [%]	66.67	Block Limit Registers [block]
Theoretical Active Warps per SM [warp]	32	Block Limit Shared Mem [block]
Achieved Occupancy [%]	65.70	Block Limit Warps [block]
Achieved Active Warps Per SM [warp]	31.54	Block Limit SM [block]

Block Limit Registers: total registers per SM / (registers per thread×threads per block)

Block Limit Shared Mem: Shared mem per SM / Shared mem per block

Block Limit Warps: Max warps per SM / Warps per block

Block Limit SM: hard architecture limit 16 blocks per SM

The register file size is 64K 32-bit registers per SM.
Warps per SM 48
<https://docs.nvidia.com/cuda/ampere-tuning-guide/index.html>

6. Launch Statistics & Occupancy

